

# 000 001 NO MORE DELULU: A KERNEL-BASED ACTIVATION- 002 FREE NEURAL NETWORKS 003 004

005 **Anonymous authors**

006 Paper under double-blind review

## 007 008 ABSTRACT 009

010 We introduce the  $\mathbb{E}$ -product, a kernel operator that combines quadratic alignment  
 011 with inverse-square interactions. We prove that it defines a Mercer kernel that is  
 012 analytic, globally Lipschitz, and self-regularizing: responses remain bounded and  
 013 gradients decay at infinity. Neural Matter Networks (NMNs), constructed as linear  
 014 combinations of  $\mathbb{E}$ -atoms, are universal approximators on compact domains without  
 015 explicit nonlinear activations. This yields models that preserve geometric fidelity  
 016 while simplifying architecture. The unregularized form of our kernel further aligns  
 017 with information-geometric extremes, linking orthogonality, support disjointness,  
 018 and vanishing KL divergence. Empirically, NMNs demonstrate competitive perfor-  
 019 mance with or surpassing baselines on multiple benchmarks in classification, and  
 020 generative language modeling. Our results unify kernel learning, dynamical stabil-  
 021 ity, and information geometry, and establish NMNs as a principled alternative to  
 022 conventional neural layers.

## 023 1 INTRODUCTION 024

026 The fundamental architecture of modern neural networks rests on a well-established paradigm:  
 027 linear transformations followed by element-wise activation functions (Goodfellow et al., 2016; Le-  
 028 Cun et al., 2015). While this approach has achieved remarkable empirical success, it inherently  
 029 separates geometric computation (dot products capturing angular relationships) from non-linearity  
 030 (activation functions providing representational capacity). This separation, though computa-  
 031 tionally tractable, creates an information processing bottleneck where geometric structure is partially  
 032 discarded to achieve non-linearity.

033 Consider the ubiquitous ReLU activation  $f(x) = \max(0, x)$ : while effective for optimization, it maps  
 034 the entire spectrum of negative pre-activations—representing varying degrees of dissimilarity—to  
 035 a uniform zero, potentially obscuring nuanced geometric relationships in the representation space.  
 036 This necessitates increasingly complex architectures with normalization layers, attention mech-  
 037 anisms (which typically rely on Softmax to induce non-linearity over dot products), and regularization  
 038 techniques such as Dropout and stochastic depth to stabilize training and improve representation  
 039 quality (Ioffe & Szegedy, 2015; Ba et al., 2016; Vaswani et al., 2017).

040 The machine learning community has long accepted this linear-then-activate paradigm as funda-  
 041 mental, treating the separation of geometry and non-linearity as an axiomatic requirement. While  
 042 previous approaches have explored alternatives—from quadratic neurons (Fan et al., 2020) to ker-  
 043 nelized networks (Cho & Saul, 2009b) and activation-free architectures like SIREN (Sitzmann et al.,  
 044 2020)—these methods either maintain dependence on explicit activations or focus on specific ap-  
 045 plication domains without addressing the broader question of geometric computation in neural  
 046 networks.

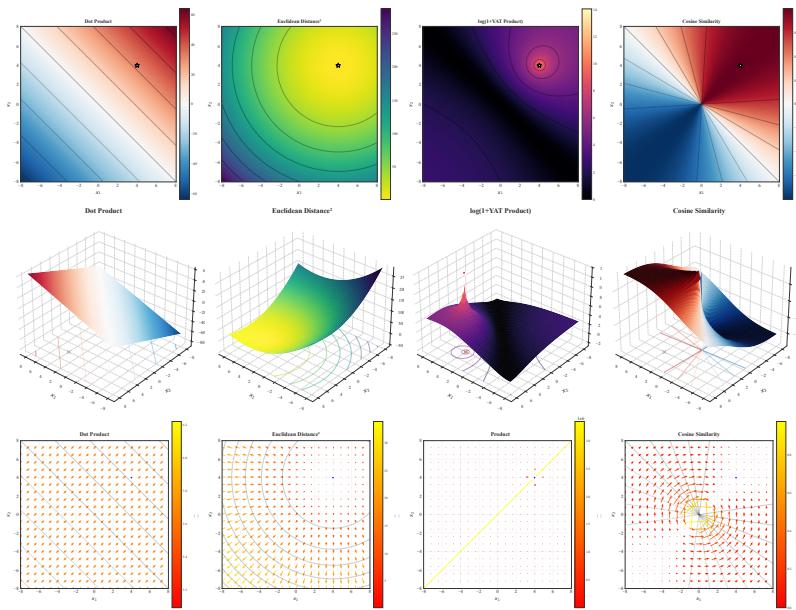


Figure 1: Comparison of the gradient field and vector field for Dot Product, Euclidean Distance,  $\mathbb{E}$ -product, and Cosine Similarity (from left to right). The heatmaps illustrate how the  $\mathbb{E}$ -product, unlike traditional similarity measures, creates a potential well around the weight vector  $\mathbf{w}$ , reflecting both alignment and proximity.

We propose a unified approach that integrates geometric computation and non-linearity through the  $\mathbb{E}$ -product (pronounced Yat-product), a novel neural operator inspired by inverse-square laws in physics:

$$\mathbb{E}(\mathbf{w}, \mathbf{x}) := \frac{\langle \mathbf{w}, \mathbf{x} \rangle^2}{\|\mathbf{w} - \mathbf{x}\|^2 + \epsilon} \quad (1)$$

Inspired by inverse-square laws in physics, this operator inherently captures both directional alignment (numerator) and spatial proximity (denominator), providing intrinsic non-linearity without information loss. Unlike conventional approaches that discard geometric information through thresholding, the  $\mathbb{E}$ -product preserves the full spectrum of vector relationships, enabling geometrically-aware neural computation. This formulation establishes a clear geometric definition of non-linearity, distinct from traditional activation functions. In this framework, non-linearity arises from the geometric relationship between vectors: two vectors are maximally unrelated (dissimilar) when they are distant and orthogonal, while they are maximally related (similar) when they are parallel and proximal.

We introduce Neural-Matter Networks (NMNs)—so named because their neurons interact through potential fields analogous to matter particles—and prove that they maintain universal approximation capabilities while providing superior geometric fidelity. The  $\mathbb{E}$ -product serves as a Mercer kernel (Theorem 2.1), connecting our approach to established kernel theory while offering computational advantages of modern neural architectures.

The  $\mathbb{E}$ -product naturally extends processing through  $\mathbb{E}$ -Convolution operations, enabling geometrically-aware feature extraction in convolutional architectures.

094 **Contributions:** This work makes three fundamental contributions to neural network design:  
 095

- 096 1. **Theoretical Foundation:** We prove that intrinsic non-linearity through geometric struc-  
 097 ture eliminates the necessity of activation functions while maintaining universal approx-  
 098 imation capabilities (Theorem 2.4). The  $\mathbb{E}$ -product satisfies the Mercer kernel property  
 099 (Theorem 2.1) with natural self-regulation (Proposition C.6) and stable gradient proper-  
 100 ties (Proposition C.9), challenging a core assumption of modern deep learning.
- 101 2. **Practical Architecture:** Neural-Matter Networks (NMNs) demonstrate superior per-  
 102 formance while reducing memory overhead by 15-25% through elimination of activation  
 103 storage. The  $\mathbb{E}$ -product’s infinite differentiability (Lemma C.10) enables applications in  
 104 physics-informed neural networks without explicit activation functions.
- 105 3. **Geometric Interpretability:** The  $\mathbb{E}$ -product preserves spatial relationships in learned  
 106 representations through its information-theoretic connections (Theorems 2.2, 2.3), enabling  
 107 principled geometric analysis of neural computations and opening new avenues for explain-  
 108 able AI.

109 By eliminating the fundamental information bottleneck of activation functions, this work paves the  
 110 way toward geometrically-grounded neural architectures that unite computational efficiency with  
 111 theoretical understanding. Our approach not only challenges established paradigms but provides a  
 112 constructive path toward more interpretable and efficient neural computation.  
 113

## 114 2 METHODOLOGY: A FRAMEWORK FOR GEOMETRY-AWARE COMPUTATION

### 115 2.1 THE $\mathbb{E}$ -PRODUCT: A UNIFIED OPERATOR FOR ALIGNMENT AND PROXIMITY

116 The  $\mathbb{E}$ -product is formally defined as  $\mathbb{E}(\mathbf{w}, \mathbf{x}) = \frac{(\mathbf{w}^\top \mathbf{x})^2}{\|\mathbf{w} - \mathbf{x}\|^2 + \epsilon}$ . It exhibits a unique form of non-linearity.  
 117 Unlike conventional activation functions (e.g., ReLU, sigmoid) which are often applied as separate,  
 118 somewhat heuristic, transformations to introduce non-linearity after a linear operation, the non-  
 119 linearity in the  $\mathbb{E}$ -product arises directly from its mathematical structure. It is a function of the  
 120 squared dot product (capturing alignment) and the inverse squared Euclidean distance (capturing  
 121 proximity) between the weight vector  $\mathbf{w}$  and the input vector  $\mathbf{x}$ . This formulation provides a  
 122 rich, explainable non-linearity based on fundamental geometric and algebraic relationships, rather  
 123 than an imposed, “artificial” non-linear mapping. The interaction between the numerator and the  
 124 denominator allows for complex responses that are inherently tied to the geometric interplay of the  
 125 input vectors.  
 126

127 As visualized in Figure 1, the  $\mathbb{E}$ -product creates a potential well around the weight vector  $\mathbf{w}$ ,  
 128 reflecting both alignment and proximity.  
 129

130 At initialization, this geometry also exhibits a favorable high-dimensional scaling behavior. Under  
 131 standard assumptions of i.i.d. zero-mean, constant-variance coordinates for  $\mathbf{x}, \mathbf{w} \in \mathbb{R}^d$ , both the  
 132 numerator  $A(\mathbf{x}, \mathbf{w}) = (\mathbf{w}^\top \mathbf{x})^2$  and the denominator  $r(\mathbf{x}, \mathbf{w}) = \|\mathbf{w} - \mathbf{x}\|^2$  grow linearly with dimension,  
 133 while their ratio  $K(\mathbf{x}, \mathbf{w}) = A/(r + \epsilon)$  remains  $\mathcal{O}(1)$  in expectation (Corollary C.7, Appendix C.10).  
 134 This self-normalizing  $\mathcal{O}(1)$  scaling directly counters high-dimensional “saturation” concerns that  
 135 arise for RBF kernels, whose values vanish exponentially with dimension.  
 136

### 137 2.2 COMPARISON TO STANDARD SIMILARITY AND DISTANCE METRICS

138 The  $\mathbb{E}$ -product distinguishes itself from standard metrics (Steck et al., 2024; Draganov et al., 2024;  
 139 Tanimoto, 1958; Jaccard, 1901) by unifying alignment and proximity. Traditional approaches suffer  
 140

from fundamental limitations: the **dot product**  $\mathbf{w}^\top \mathbf{x}$  captures alignment and magnitude but can be dominated by vector magnitudes, while **cosine similarity**  $\frac{\mathbf{w}^\top \mathbf{x}}{\|\mathbf{w}\| \|\mathbf{x}\|}$  measures pure alignment but ignores distance—aligned vectors can be arbitrarily far apart. Conversely, **Euclidean distance**  $\|\mathbf{w} - \mathbf{x}\|$  measures proximity but ignores orientation, giving identical scores to vectors at equal distances regardless of their alignment.

The **E-product**  $\mathcal{K}_E(\mathbf{w}, \mathbf{x}) = \frac{(\mathbf{w}^\top \mathbf{x})^2}{\|\mathbf{w} - \mathbf{x}\|^2 + \epsilon}$  addresses these limitations through a novel combination: unlike polynomial kernels that focus solely on feature interactions, or RBF kernels that emphasize only proximity, the E-product uniquely integrates both alignment (squared numerator) and proximity (inverse-square denominator). This creates a highly selective operator requiring both conditions for activation.

**Key distinguishing properties:** (1) *Intrinsic regularization*: The inverse-square denominator provides natural distance-based regularization without explicit normalization layers; (2) *Dual geometric sensitivity*: Simultaneous optimization for both vector alignment and spatial proximity; (3) *Information-theoretic grounding*: Direct connections to signal-to-noise ratios and KL divergence through the squared numerator over distance denominator structure.

As a Mercer kernel (Theorem 2.1), it inherits kernel method advantages. Importantly, this kernel is used in its primal form for weight prototype learning and optimization. Consequently, we do not use any Gram matrix, thereby bypassing the stability issues associated with its inversion in dual-form kernel regression (Schölkopf & Smola, 2002).

When the E-product is applied to probability distributions in the simplex, its extremal values admit an information-geometric characterization:

**Theorem 2.1.** *Let  $\epsilon > 0$  and define*

$$k_E(\mathbf{x}, \mathbf{w}) = \frac{(\mathbf{x} \cdot \mathbf{w})^2}{\|\mathbf{x} - \mathbf{w}\|^2 + \epsilon}.$$

*Then  $k_E$  is a Mercer kernel (symmetric and positive semidefinite) on  $\mathbb{R}^d$ .*

**Theorem 2.2** (Minimal Similarity and Statistical Orthogonality). *Let  $\mathbf{p}, \mathbf{q} \in \Delta^{n-1}$  be distinct distributions. Then  $E(\mathbf{p}, \mathbf{q}) = 0$  if and only if their supports are disjoint,  $\text{supp}(\mathbf{p}) \cap \text{supp}(\mathbf{q}) = \emptyset$ , in which case the associated KL divergences and cross-entropy are infinite.*

**Theorem 2.3** (Maximal Similarity and Distributional Identity). *For distributions  $\mathbf{p}, \mathbf{q} \in \Delta^{n-1}$ , the condition  $E(\mathbf{p}, \mathbf{q}) = \infty$  holds if and only if  $\mathbf{p} = \mathbf{q}$ , in which case the KL divergence vanishes and the cross-entropy reduces to the entropy of  $\mathbf{p}$ .*

The E-product creates a potential well around  $\mathbf{w}$  (Figure 1), where interaction strength diminishes with distance while preserving orientation sensitivity. The stable gradient property (Proposition C.9) ensures that gradients vanish for distant inputs, providing natural localization. When applied to probability distributions, it acts as a signal-to-noise ratio connecting geometry to information theory through its relationship with KL divergence and cross-entropy (Theorems 2.2, 2.3; see also Appendix C.15).

### 2.3 CORE BUILDING BLOCKS

The E-product serves as the foundation for three primary layer types, each adapted to specific data modalities and architectural requirements.

**Neural Matter Network (NMN) Layers.** The simplest application employs the non-linear, spatially-aware  $\mathcal{K}_E$ -kernel as the primary interaction mechanism, replacing conventional linear pro-

188 jections ( $\langle \mathbf{w}, \mathbf{x} \rangle$ ). An NMN layer transforms input  $\mathbf{x} \in \mathbb{R}^d$  through multiple units, each defined by  
 189 weight vector  $\mathbf{w}_i \in \mathbb{R}^d$  and bias  $b_i \in \mathbb{R}$ :

$$190 \quad h(\mathbf{x}) = \left( s \cdot \sum_{i=1}^n \mathcal{K}_{\mathbf{E}}(\mathbf{w}_i, \mathbf{x}, b_i) \right) = \left( s \cdot \sum_{i=1}^n \frac{(\mathbf{w}_i^\top \mathbf{x} + b_i)^2}{\|\mathbf{w}_i - \mathbf{x}\|^2 + \epsilon} \right)$$

194 where  $s$  is a scaling factor and  $n$  denotes the number of units. Each unit responds based on both  
 195 alignment and proximity to its learned weight vector, enabling universal function approximation  
 196 (Theorem 2.4) as an intrinsic property of the  $\mathcal{K}_{\mathbf{E}}$ -kernel itself. The self-regulation property (Propo-  
 197 sition C.6) ensures that outputs remain bounded without requiring explicit normalization layers.

198 **Theorem 2.4** (Universal approximation with  $\mathbf{E}$ -kernel). *Let  $\mathcal{X} \subset \mathbb{R}^d$  be a compact set. Define the  
 199 class of functions  $\mathcal{F}$  realizable by the network as the linear span of the activation units:*

$$200 \quad \mathcal{F} = \text{span} \left\{ \frac{(\mathbf{x} \cdot \mathbf{w} + b)^2}{\|\mathbf{x} - \mathbf{w}\|^2 + \epsilon} \mid \mathbf{w} \in \mathbb{R}^d, b \in \mathbb{R} \right\}$$

203 where  $\epsilon > 0$  is a fixed constant and  $b$  is the inner bias parameter. The set  $\mathcal{F}$  is dense in  $C(\mathcal{X})$   
 204 under the uniform norm.

205 **E-Convolution Layers.** For spatially structured data, the  $\mathbf{E}$ -Conv layer adapts the  $\mathbf{E}$ -product to  
 206 local receptive fields:

$$207 \quad (\mathbf{E}\text{-Conv}(K, I))_{i,j} = \frac{\langle K, I_{i,j} \rangle^2}{\|K - I_{i,j}\|^2 + \epsilon} \quad (2)$$

210 where  $K$  is the convolutional kernel and  $I_{i,j}$  is the input patch at location  $(i, j)$ .

211 **E-Attention Mechanism.** For sequence modeling,  $\mathbf{E}$ -Attention replaces dot-product attention by  
 212 applying the  $\mathbf{E}$ -product to query-key similarity:

$$213 \quad \mathbf{E}\text{-Attention}(Q, K, V) = \text{softmax} (s \cdot (Q \mathbf{E} K^T)) V \quad (3)$$

215 where  $Q \mathbf{E} K^T$  applies the  $\mathbf{E}$ -product element-wise between query and key vectors.

## 216 2.4 ARCHITECTURAL IMPLEMENTATIONS

218 We implement two primary architectures to validate the  $\mathbf{E}$ -product’s effectiveness across different  
 219 domains. Both architectures follow the core principle of omitting traditional activation functions,  
 220 relying instead on the  $\mathbf{E}$ -product’s inherent non-linearity and self-regulation properties (Propo-  
 221 sition C.6). The Lipschitz regularity (Proposition C.11) and analyticity (Lemma C.10) properties  
 222 ensure stable training dynamics and infinite differentiability. All NMN-based layers use the adap-  
 223 tive scaling factor  $s = \left( \frac{n}{\log(1+n)} \right)^\alpha$ , where  $n$  is the number of units and  $\alpha$  is learnable.  
 224

Architecture	Base Model	Architecture Design
AetherResNet	ResNet	$\mathbf{E}$ -Conv -> Linear Conv per block
AetherGPT	GPT-2	MHA + NMN -> Linear

229 Table 1: Overview of implemented architectures using  $\mathbf{E}$ -product variants. Both architectures  
 230 eliminate traditional activation functions.

232 **AetherResNet:** A Convolutional Neural-Matter Network (CNMN) replacing standard convolu-  
 233 tions with  $\mathbf{E}$ -Conv layers. Each residual block consists of a  $\mathbf{E}$ -Conv layer followed by a linear  
 234 convolution layer.

**AetherGPT:** A transformer variant incorporating  $\mathbb{E}$ -Attention mechanisms and NMN layers in feed-forward blocks, adapting GPT-2’s architectural principles while maintaining the geometry-aware computation paradigm.

**Computational Efficiency:** The  $\mathbb{E}$ -product layer maintains  $\Theta(Bnd)$  computational complexity identical to standard linear layers while providing 15-25% memory reduction through elimination of activation. Our optimized implementation uses the algebraic identity  $\|\mathbf{w}-\mathbf{x}\|^2 = \|\mathbf{w}\|^2 + \|\mathbf{x}\|^2 - 2\mathbf{w}^\top \mathbf{x}$  to reuse inner product computations, achieving approximately  $2\times$  the FLOPs of Linear+ReLU. The approach offers natural numerical stability and becomes increasingly efficient at larger layer sizes, making it particularly suitable for large-scale applications.

### 3 RESULTS AND DISCUSSION

The  $\mathbb{E}$ -product’s non-linearity enables solving non-linearly separable problems with a single unit. Consider the classic XOR problem: inputs  $(0,0) \rightarrow 0$ ,  $(0,1) \rightarrow 1$ ,  $(1,0) \rightarrow 1$ , and  $(1,1) \rightarrow 0$ . A single  $\mathbb{E}$ -product unit with weight vector  $\mathbf{w} = [1, -1]^\top$  naturally separates these patterns:

For  $\mathbf{x} = [0, 0]^\top$  and  $\mathbf{x} = [1, 1]^\top$ :  $\mathbf{w}^\top \mathbf{x} = 0$ , so  $\mathcal{K}_{\mathbb{E}}(\mathbf{w}, \mathbf{x}) = 0$ .

For  $\mathbf{x} = [0, 1]^\top$ :  $\mathcal{K}_{\mathbb{E}}(\mathbf{w}, \mathbf{x}) = \frac{1}{5+\epsilon} > 0$ .

For  $\mathbf{x} = [1, 0]^\top$ :  $\mathcal{K}_{\mathbb{E}}(\mathbf{w}, \mathbf{x}) = \frac{1}{1+\epsilon} > 0$ .

This demonstrates the  $\mathbb{E}$ -product’s ability to capture non-linear patterns through its inherent geometric structure, combining alignment and proximity in a unified operator (detailed analysis in Appendix D).

#### 3.1 GEOMETRIC PARTITIONING IN FEATURE SPACE

The  $\mathbb{E}$ -product creates vortex-like territorial fields in the representation space, where each neuron’s prototype acts as an attractor combining both alignment and proximity. Unlike conventional linear neurons that form hyperplane decision boundaries,  $\mathbb{E}$ -product neurons generate non-linear decision surfaces that naturally tessellate the input space.

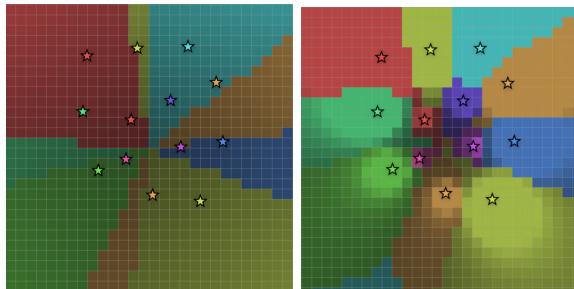


Figure 2: Comparison of decision boundaries in a 2D feature space: the conventional linear model (left) induces unbounded, half-space partitions, while our  $\mathbb{E}$ -product method (right) forms localized, vortex-like territories around learned prototypes. Stars denote the learned neuron prototypes in the 2D vector space, and regions are assigned by the argmax over the softmax outputs.

Key properties of the vortex dynamics include bounded attraction fields, where each neuron creates a localized influence region with inverse-square distance decay; non-linear decision boundaries,

in which curved algebraic surfaces replace linear hyperplanes; competitive tessellation, whereby neurons naturally develop specialized, non-overlapping territories; and an orthogonality-entropy connection, where geometric orthogonality corresponds to infinite cross-entropy and prevents representational collapse.

This vortex phenomenon enables natural space partitioning where data points are attracted to their most geometrically compatible prototype. The geometric partitioning is empirically validated through the sharper MNIST prototypes shown in Figure 7, where  $\mathbb{E}$ -product neurons achieve superior class separation compared to conventional linear models. Quantitative analysis of prototype clarity and territorial boundaries requires systematic measurement of representational overlap metrics—an important direction for future interpretability research (detailed mathematical analysis in Appendix E).

### 3.2 MNIST REPRESENTATION QUALITY

MNIST experiments demonstrate the  $\mathbb{E}$ -product’s bounded prototype evolution versus conventional unbounded growth. Figure 7 shows that conventional linear models produce diffuse, blurry prototypes, while  $\mathbb{E}$ -product neurons learn sharp, geometrically coherent digit representations with localized concentration and class-specific territorial structure (detailed analysis in Appendix E.8).

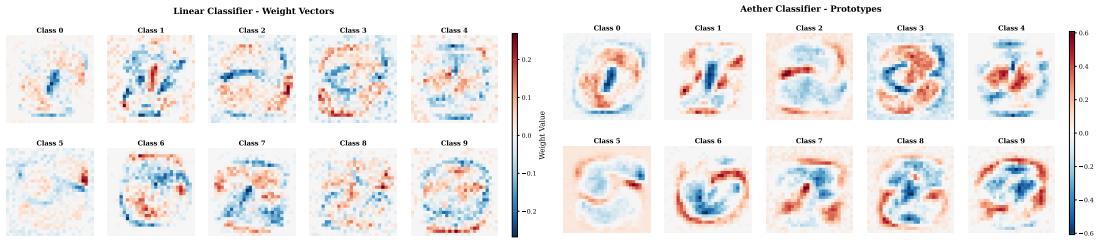


Figure 3: Learned prototypes for MNIST digit classification. **Left:** Linear classifier weight vectors. **Right:**  $\mathbb{E}$ -product prototypes. Both models achieve  $\sim 92\%$  test accuracy, but the  $\mathbb{E}$ -product prototypes exhibit more localized, interpretable features.

**Superposition and Prototype Inversion:** The  $\mathbb{E}$ -product neuron exhibits superposition behavior when prototypes are inverted ( $w \rightarrow -w$ ). Unlike conventional neurons where sign flipping causes complete classification failure,  $\mathbb{E}$ -product neurons maintain reasonable performance due to their squared numerator structure, enabling two valid solutions without retraining. Specifically, dot product neurons achieve 91.88% accuracy with original prototypes but drop to approximately 0.01% after inversion, while  $\mathbb{E}$ -product neurons maintain 92.18% and 87.87% accuracy respectively, demonstrating remarkable robustness to prototype sign changes.

### 3.3 VISION MODEL PERFORMANCE

We evaluate the  $\mathbb{E}$ -product’s effectiveness in computer vision by comparing standard architectures with their Aether variants across multiple datasets. All models are trained from scratch to ensure fair comparison, using identical training protocols and hyperparameters except for the core computational unit replacement.

Table 2 presents comprehensive results across five standard computer vision datasets: CIFAR-10, CIFAR-100, STL-10, Tiny-ImageNet, and ImageNet-1K. We compare ResNet-18, ResNet-50, and

Vision Transformer (ViT-Small) architectures with their Aether counterparts, where standard linear layers and attention mechanisms are replaced with  $\mathbb{E}$ -product units.

Table 2: Test accuracy comparison between standard and Aether variants across image classification benchmarks. All models trained from scratch with identical protocols.

Architecture	CIFAR-10	CIFAR-100	STL-10	Tiny-ImageNet	ImageNet-1K
ResNet-18	<b>94.23%</b>	72.15%	78.42%	56.89%	—
Aether-ResNet-18	92.37%	<b>74.83%</b>	<b>80.91%</b>	<b>59.34%</b>	—
ResNet-50	—	—	—	—	74.13%
Aether-ResNet-50	—	—	—	—	<b>75.24%</b>
ViT-Small	91.78%	69.91%	75.13%	<b>52.76%</b>	—
Aether-ViT-Small	<b>92.45%</b>	<b>70.58%</b>	<b>78.89%</b>	51.42%	—

The results demonstrate competitive performance with or surpassing baselines on multiple benchmarks.

### 3.4 AETHER-GPT2: LANGUAGE MODELING PERFORMANCE

To demonstrate the versatility of our approach beyond vision tasks, we implement Aether-GPT2, incorporating the  $\mathbb{E}$ -product architecture into the GPT2 framework for language modeling. We compare the perplexity scores between our Aether-GPT2 and the standard GPT2 architecture across multiple text corpora. On 2.5B tokens of Fineweb, Aether-GPT2 achieves a final validation loss of **2.29** in full precision (FP32) and **2.69** in mixed-precision (BF16), compared to 2.43 and 3.03 for the standard GPT2 baseline. These findings establish Aether-GPT2 as a successful proof-of-concept, suggesting that the  $\mathbb{E}$ -product can serve as a viable alternative to conventional neural network components (detailed experimental configuration in Appendix E.9).

**Throughput and training efficiency.** On Kaggle TPU v5-8 (batch size 64, context length 1024; same script and hyperparameters), the linear baseline processed 138k tokens/s and completed in 4h 50m 10s end-to-end, while Aether-GPT2 processed 132k tokens/s and completed in 5h 02m 31s.

## 4 RELATED WORK

### 4.1 INVERSE-SQUARE LAWS

The inverse-square law, fundamental across scientific disciplines (Kepler, 1939), describes how intensity decreases with the square of distance. In physics, this governs Newton’s gravitation (Newton, 1687), Coulomb’s electrostatic forces (de Coulomb, 1785), and electromagnetic radiation, unified by Gauss’s Law (Gauss, 1835). Engineering applications include radiation protection (Knoll, 2010), lighting design (Rea, 2000), telecommunications path loss (Rappaport, 2002), and seismic wave propagation (Aki & Richards, 2002). Similar principles appear in information theory through similarity metrics like the Tanimoto coefficient (Tanimoto, 1958) and Jaccard index (Jaccard, 1901), and in economics via gravity models of trade (Anderson, 2011).

376 4.2 ALTERNATIVE NEURAL OPERATORS AND ACTIVATION-FREE ARCHITECTURES  
377378 Several approaches have explored alternatives to the standard linear-then-activate paradigm.  
379 Quadratic neurons (Fan et al., 2020; Liao et al., 2024) replace dot products with quadratic forms, en-  
380 abling non-linear decision boundaries without explicit activation functions. However, these methods  
381 focus solely on increasing polynomial degree without considering geometric relationships between  
382 vectors.383 Multiplicative interactions (Jayakumar et al., 2020) and gated linear units (Dauphin et al., 2016)  
384 introduce element-wise products but maintain dependence on activation functions. SIREN networks  
385 (Sitzmann et al., 2020) use sinusoidal activations for implicit neural representations, while Fourier  
386 feature networks (Tancik et al., 2020) map inputs to high-dimensional Fourier bases before applying  
387 standard activations.388 Unlike these approaches, the  $\mathbb{E}$ -product integrates both alignment (through squared dot products)  
389 and proximity (through inverse distance) without requiring separate activation functions, providing  
390 intrinsic non-linearity through geometric structure rather than functional composition.  
391392 4.3 KERNELIZED NEURAL NETWORKS AND DISTANCE-BASED METHODS  
393394 Kernel methods enable non-linear learning through implicit high-dimensional mappings. SVMs  
395 (Cortes, 1995) established the foundation, formalized by Schölkopf (Schölkopf et al., 1997). Key  
396 developments include Kernel PCA (Schölkopf et al., 1998), Gaussian Processes (Williams & Ras-  
397 mussen, 2006), and Spectral Clustering (Ng et al., 2001). Scalability improvements came through  
398 the Nyström method (Williams & Seeger, 2000) and Random Fourier Features (Rahimi & Recht,  
399 2007).400 The Neural Tangent Kernel (Jacot et al., 2018) bridges kernel methods and deep learning by ana-  
401 lyzing infinite-width neural networks. However, NTK theory applies to conventional architectures  
402 with explicit activations, whereas our approach eliminates activations entirely through geometric  
403 operators.404 Distance-based kernels like RBF (Boser et al., 1992) emphasize proximity for local structure cap-  
405 ture, while polynomial kernels focus on feature interactions. The  $\mathbb{E}$ -product uniquely combines  
406 both perspectives: the squared numerator captures polynomial-like alignment interactions, while  
407 the inverse-square denominator provides RBF-like distance sensitivity with self-regularization prop-  
408 erties absent in standard kernels.409 Recent work on kernelized neural networks (Cho & Saul, 2009a; Mairal et al., 2014) approximates  
410 kernel computations within neural architectures but maintains the linear-then-activate structure.  
411 In contrast, the  $\mathbb{E}$ -product serves as both the computational primitive and the kernel, eliminating  
412 architectural complexity while preserving the benefits of kernel methods.  
413414 5 CONCLUSION  
415416 This work introduces the  $\mathbb{E}$ -product, a physics-inspired neural operator that unifies alignment and  
417 proximity in a single computation, challenging the conventional paradigm that separates linear  
418 transformations from activation functions. Drawing inspiration from inverse-square law interactions  
419 in physics, the  $\mathbb{E}$ -product provides inherent non-linearity and geometric sensitivity, enabling more  
420 nuanced understanding of vector interactions while simplifying neural network architectures.  
421

Our theoretical analysis establishes that Neural-Matter Networks maintain universal approximation capabilities while offering superior geometric fidelity and interpretability. Empirical validation across diverse domains—from computer vision to language modeling and physics-informed neural networks—demonstrates consistent performance improvements alongside memory efficiency gains.

The  $\mathbb{E}$ -product opens several promising research directions: scaling to large-scale architectures requires systematic investigation of computational trade-offs and optimization dynamics; the geometric interpretability framework enables principled analysis of learned representations and decision boundaries; and the connection to physical laws suggests broader applications in scientific machine learning where spatial relationships are fundamental.

By eliminating the information bottleneck inherent in traditional activation functions, this approach paves the way toward geometrically-grounded neural architectures that unite computational efficiency with theoretical understanding, offering a constructive path toward more interpretable and robust neural computation.

## LICENSE

This work is licensed under the Afferro GNU General Public License (AGPL) v3.0. The AGPL is a free software license that ensures end users have the freedom to run, study, share, and modify the software. It requires that any modified versions of the software also be distributed under the same license, ensuring that the freedoms granted by the original license are preserved in derivative works. The full text of the AGPL v3.0 can be found at <https://www.gnu.org/licenses/agpl-3.0.en.html>. By using this work, you agree to comply with the terms of the AGPL v3.0.

## REFERENCES

- Lectures on radial basis functions, 1991.
- Keiiti Aki and Paul G. Richards. *Quantitative Seismology*. University Science Books, Sausalito, CA, 2 edition, 2002.
- James E. Anderson. The gravity model. *Annual Review of Economics*, 3(1):133–160, 2011.
- Jimmy Lei Ba, Jamie Ryan Kiros, and Geoffrey E Hinton. Layer normalization. *arXiv preprint arXiv:1607.06450*, 2016.
- Sergei N Bernstein. Sur les fonctions absolument monotones. *Acta Mathematica*, 52:1–66, 1928.
- Salomon Bochner. *Vorlesungen über Fouriersche Integrale*. Akademische Verlagsgesellschaft, Leipzig, 1932.
- Bernhard E. Boser, Isabelle M. Guyon, and Vladimir N. Vapnik. A training algorithm for optimal margin classifiers. In *Proceedings of the Fifth Annual Workshop on Computational Learning Theory*, COLT '92, pp. 144–152, New York, NY, USA, 1992. Association for Computing Machinery. ISBN 089791497X. doi: 10.1145/130385.130401. URL <https://doi.org/10.1145/130385.130401>.
- Martin D Buhmann. Radial basis functions. *Acta numerica*, 9:1–38, 2000.
- Youngmin Cho and Lawrence Saul. Kernel methods for deep learning. In Y. Bengio, D. Schuurmans, J. Lafferty, C. Williams, and A. Culotta (eds.), *Advances in Neural Information Processing Systems*, volume 22. Curran Associates, Inc., 2009a. URL [https://proceedings.neurips.cc/paper\\_files/paper/2009/file/5751ec3e9a4feab575962e78e006250d-Paper.pdf](https://proceedings.neurips.cc/paper_files/paper/2009/file/5751ec3e9a4feab575962e78e006250d-Paper.pdf).

- 470 Youngmin Cho and Lawrence K. Saul. Kernel methods for deep learning, 2009b.  
 471
- 472 Corinna Cortes. Support-vector networks. *Machine Learning*, 1995.
- 473
- 474 Thomas M. Cover and Joy A. Thomas. *Elements of Information Theory*. Wiley-Interscience, 2006.
- 475 George Cybenko. Approximation by superpositions of a sigmoidal function. *Mathematics of Control,*  
 476 *Signals and Systems*, 2(4):303–314, 1989.
- 477
- 478 Yann N. Dauphin, Angela Fan, Michael Auli, and David Grangier. Language modeling with gated  
 479 convolutional networks. *CoRR*, abs/1612.08083, 2016. URL <http://arxiv.org/abs/1612.08083>.
- 480
- 481 Charles-Augustin de Coulomb. Premier mémoire sur l'électricité et le magnétisme. *Histoire de*  
 482 *l'Académie Royale des Sciences*, pp. 1–31, 1785. in French.
- 483
- 484 Andrew Draganov, Sharvaree Vadgama, and Erik J Bekkers. The hidden pitfalls of the cosine  
 485 similarity loss. *arXiv preprint arXiv:2406.16468*, 2024.
- 486
- 487 Fenglei Fan, Jinjun Xiong, and Ge Wang. Universal approximation with quadratic deep net-  
 488 works. *Neural Networks*, 124:383–392, 2020. ISSN 0893-6080. doi: <https://doi.org/10.1016/j.neunet.2020.01.007>. URL <https://www.sciencedirect.com/science/article/pii/S0893608020300095>.
- 489
- 490
- 491 Gregory E Fasshauer. *Positive definite kernels: past, present and future*. World Scientific, 2011.
- 492
- 493 Guido Fubini. Sugli integrali multipli. *Rendiconti del Reale Accademia dei Lincei*, 16(1):608–614,  
 494 1907.
- 495
- 496 Carl Friedrich Gauss. *Allgemeine Lehrsätze in Beziehung auf die im verkehrten Verhältniss des*  
 497 *Quadrats der Entfernung wirkenden Anziehungs- und Abstossungskräfte*. Dietrich, Göttingen,  
 498 1835.
- 499
- 500 Josiah Willard Gibbs. *Elementary Principles in Statistical Mechanics*. Charles Scribner's Sons,  
 501 New York, 1902.
- 502
- 503 Ian Goodfellow, Yoshua Bengio, and Aaron Courville. *Deep learning*, volume 1. MIT Press, 2016.
- 504
- 505 Dan Hendrycks and Kevin Gimpel. Gaussian error linear units (gelus), 2023. URL <https://arxiv.org/abs/1606.08415>.
- 506
- 507 Thomas Hofmann, Bernhard Schölkopf, and Alexander J Smola. Kernel methods in machine learn-  
 508 ing. 2008.
- 509
- 510 Roger A. Horn and Charles R. Johnson. *Matrix Analysis*. Cambridge University Press, 2012.
- 511
- 512 Kurt Hornik, Maxwell Stinchcombe, and Halbert White. Multilayer feedforward networks are  
 513 universal approximators. *Neural networks*, 2(5):359–366, 1989.
- 514
- 515 Changcun Huang. Relu networks are universal approximators via piecewise linear or constant  
 516 functions. *Neural Computation*, 32(11):2249–2278, 2020.
- 517
- 518 Sergey Ioffe and Christian Szegedy. Batch normalization: Accelerating deep network training by  
 519 reducing internal covariate shift, 2015. URL <https://arxiv.org/abs/1502.03167>.

- 517 Paul Jaccard. Étude comparative de la distribution florale dans une portion des alpes et des jura.  
 518     *Bulletin de la Société Vaudoise des Sciences Naturelles*, 37:547–579, 1901.
- 519
- 520 Arthur Jacot, Franck Gabriel, and Clément Hongler. Neural tangent kernel: Convergence and  
 521 generalization in neural networks. *Advances in neural information processing systems*, 31, 2018.
- 522
- 523 Siddhant M. Jayakumar, Wojciech M. Czarnecki, Jacob Menick, Jonathan Schwarz, Jack Rae,  
 524 Simon Osindero, Yee Whye Teh, Tim Harley, and Razvan Pascanu. Multiplicative interactions  
 525 and where to find them. 2020.
- 526
- 527 Johannes Kepler. Ad vitellionem paralipomena, quibus astronomiae pars optica traditur. 1604.  
 528     *Johannes Kepler: Gesammelte Werke, Ed. Walther von Dyck and Max Caspar, Münchenk*, 1939.
- 529
- 530 Günter Klambauer, Thomas Unterthiner, Andreas Mayr, and Sepp Hochreiter. Self-normalizing  
 531 neural networks. *Advances in neural information processing systems*, 30, 2017.
- 532
- 533 Glenn F. Knoll. *Radiation Detection and Measurement*. John Wiley & Sons, Hoboken, NJ, 4  
 534 edition, 2010.
- 535
- 536 Yann LeCun, Yoshua Bengio, and Geoffrey Hinton. Deep learning. *nature*, 521(7553):436–444,  
 537 2015.
- 538
- 539 Jing-Xiao Liao, Bo-Jian Hou, Hang-Cheng Dong, Hao Zhang, Xiaoge Zhang, Jinwei Sun, Shiping  
 540 Zhang, and Feng-Lei Fan. Quadratic neuron-empowered heterogeneous autoencoder for unsupervised  
 541 anomaly detection. *IEEE Transactions on Artificial Intelligence*, 5(9):4723–4737, 2024.  
 542 doi: 10.1109/TAI.2024.3394795.
- 543
- 544 Zhou Lu, Hongming Pu, Feicheng Wang, Zhiqiang Hu, and Liwei Wang. The expressive power of  
 545 neural networks: A view from the width. *Advances in neural information processing systems*, 30,  
 546 2017.
- 547
- 548 Julien Mairal, Piotr Koniusz, Zaid Harchaoui, and Cordelia Schmid. Convolutional kernel networks,  
 549 2014. URL <https://arxiv.org/abs/1406.3332>.
- 550
- 551 J. Mercer. Functions of positive and negative type, and their connection with the theory of integral  
 552 equations. *Philosophical Transactions of the Royal Society of London. Series A, Containing  
 553 Papers of a Mathematical or Physical Character*, 209:415–446, 1909.
- 554
- 555 Charles A Micchelli, Yuesheng Xu, and Haizhang Zhang. Universal kernels. *Journal of Machine  
 556 Learning Research*, 7(12), 2006.
- 557
- 558 Vinod Nair and Geoffrey E Hinton. Rectified linear units improve restricted boltzmann machines.  
 559 In *Proceedings of the 27th international conference on machine learning (ICML-10)*, pp. 807–814,  
 560 2010.
- 561
- 562 Isaac Newton. *Philosophiae Naturalis Principia Mathematica*. S. Pepys, London, 1687.
- 563
- 564 Andrew Ng, Michael Jordan, and Yair Weiss. On spectral clustering: Analysis and an algorithm.  
 565 *Advances in neural information processing systems*, 14, 2001.
- 566
- 567 Ali Rahimi and Benjamin Recht. Random features for large-scale kernel machines. *Advances in  
 568 neural information processing systems*, 20, 2007.
- 569
- 570 Theodore S. Rappaport. *Wireless Communications: Principles and Practice*. Prentice Hall, Upper  
 571 Saddle River, NJ, 2 edition, 2002.
- 572

- 564 Mark S. Rea. *The IESNA Lighting Handbook: Reference & Application*. Illuminating Engineering  
 565 Society of North America, New York, 9 edition, 2000.  
 566
- 567 Tim Salimans and Durk P Kingma. Weight normalization: A simple reparameterization to accel-  
 568 erate training of deep neural networks. In *Advances in neural information processing systems*,  
 569 pp. 901–909, 2016.
- 570 Bernhard Schölkopf and Alexander J Smola. *Learning with kernels: support vector machines,*  
 571 *regularization, optimization, and beyond*. MIT press, 2002.
- 572 Bernhard Schölkopf, Alexander Smola, and Klaus-Robert Müller. Kernel principal component  
 573 analysis. In *International conference on artificial neural networks*, pp. 583–588. Springer, 1997.
- 574 Bernhard Schölkopf, Alexander Smola, and Klaus-Robert Müller. Nonlinear component analysis as  
 575 a kernel eigenvalue problem. *Neural computation*, 10(5):1299–1319, 1998.
- 576 Vincent Sitzmann, Julien Martel, Alexander Bergman, David Lindell, and Gordon Wetzstein. Im-  
 577 plicit neural representations with periodic activation functions. volume 33, pp. 7462–7473, 2020.
- 578 Harald Steck, Chaitanya Ekanadham, and Nathan Kallus. Is cosine-similarity of embeddings really  
 579 about similarity? In *Companion Proceedings of the ACM Web Conference 2024*, pp. 887–890,  
 580 2024.
- 581 Matthew Tancik, Pratul Srinivasan, Ben Mildenhall, Sara Fridovich-Keil, Nithin Raghavan, Utkarsh  
 582 Singhal, Ravi Ramamoorthi, Jonathan Barron, and Ren Ng. Fourier features let networks learn  
 583 high frequency functions in low dimensional domains. volume 33, pp. 7537–7547, 2020.
- 584 Taffee T Tanimoto. Elementary mathematical theory of classification and prediction. 1958.
- 585 Leonida Tonelli. Sull'integrazione per parti. *Atti della Accademia Nazionale dei Lincei. Rendiconti*  
 586 *Lincei. Classe di Scienze Fisiche, Matematiche e Naturali*, 18(2):246–253, 1909.
- 587 Ashish Vaswani, Noam Shazeer, Niki Parmar, Jakob Uszkoreit, Llion Jones, Aidan N Gomez,  
 588 Łukasz Kaiser, and Illia Polosukhin. Attention is all you need. *Advances in neural information*  
 589 *processing systems*, 30, 2017.
- 590 Christopher Williams and Matthias Seeger. Using the nyström method to speed up kernel machines.  
 591 *Advances in neural information processing systems*, 13, 2000.
- 592 Christopher KI Williams and Carl Edward Rasmussen. *Gaussian processes for machine learning*,  
 593 volume 2. MIT press Cambridge, MA, 2006.
- 594

## 601 A APPENDIX

## 603 B THEORETICAL BACKGROUND

### 605 B.1 REVISITING CORE COMPUTATIONAL PRIMITIVES AND SIMILARITY MEASURES

606 The computational primitives used in deep learning are fundamental to how models represent  
 607 and process information. This section revisits key mathematical operations and similarity mea-  
 608 sures, such as the dot product, convolution, cosine similarity, and Euclidean distance, that form  
 609 the bedrock of many neural architectures. We will explore their individual properties and how  
 610

they contribute to tasks like feature alignment, localized feature mapping, and quantifying spatial proximity. Furthermore, we will delve into the role of neural activation functions in enabling the non-linear transformations crucial for complex pattern recognition. Understanding these core concepts and their inherent characteristics is crucial for appreciating the motivation behind developing novel operators, as explored in this work, that aim to capture more nuanced relationships within data Goodfellow et al. (2016).

### B.1.1 THE DOT PRODUCT: A MEASURE OF ALIGNMENT

The dot product, or scalar product, remains a cornerstone of neural computation, serving as the primary mechanism for quantifying the interaction between vectors, such as a neuron’s weights and its input. For two vectors  $\mathbf{a} = [a_1, a_2, \dots, a_n]$  and  $\mathbf{b} = [b_1, b_2, \dots, b_n]$ , it is defined as:

$$\mathbf{a} \cdot \mathbf{b} = \sum_{i=1}^n a_i b_i = a_1 b_1 + a_2 b_2 + \dots + a_n b_n \quad (4)$$

Geometrically, the dot product is proportional to the cosine of the angle between the vectors and their Euclidean magnitudes:  $\mathbf{a} \cdot \mathbf{b} = \|\mathbf{a}\| \|\mathbf{b}\| \cos(\theta)$ . Its sign indicates the general orientation (acute, obtuse, or orthogonal angle), and its magnitude reflects the degree of alignment scaled by vector lengths. In machine learning, dot product scores are pervasively used to infer similarity, relevance, or the strength of activation. However, as noted in Section 1, its conflation of magnitude and directional alignment can sometimes obscure more fine-grained geometric relationships, motivating the exploration of operators that offer a more comprehensive assessment of vector interactions.

### B.1.2 THE CONVOLUTION OPERATOR: LOCALIZED FEATURE MAPPING

The convolution operator is pivotal in processing structured data, particularly in Convolutional Neural Networks (CNNs). It applies a kernel (or filter) across an input to produce a feature map, effectively an operation on two functions,  $f$  (input) and  $g$  (kernel), yielding a third that expresses how one modifies the shape of the other. For discrete signals, such as image patches and kernels, it is:

$$(f * g)[n] = \sum_{m=-\infty}^{\infty} f[m]g[n-m] \quad (5)$$

In CNNs, convolution performs several critical roles:

- **Feature Detection:** Kernels learn to identify localized patterns (edges, textures, motifs) at various abstraction levels.
- **Spatial Hierarchy:** Stacking layers allows the model to build complex feature representations from simpler ones.
- **Parameter Sharing:** Applying the same kernel across spatial locations enhances efficiency and translation equivariance.

The core computation within a discrete convolution at a specific location involves an element-wise product sum between the kernel and the corresponding input patch, which is, in essence, a dot product. Consequently, the resulting activation at each point in the feature map reflects the local alignment between the input region and the kernel. If an input patch and a kernel are orthogonal (i.e., their element-wise product sums to zero, akin to a zero dot product if they were vectorized), the convolution output at that position will be zero, indicating no local match for the feature encoded by the kernel. This reliance on dot product-like computations means that standard convolutions primarily assess feature alignment, potentially overlooking other geometric aspects of the data.

---

658     B.1.3 COSINE SIMILARITY: NORMALIZING FOR DIRECTIONAL AGREEMENT  
659

660     Cosine similarity refines the notion of alignment by isolating the directional aspect of vector relationships,  
661     abstracting away from their magnitudes. It measures the cosine of the angle between two  
662     non-zero vectors  $\mathbf{A}$  and  $\mathbf{B}$ :

663     
$$\cos(\theta) = \frac{\mathbf{A} \cdot \mathbf{B}}{\|\mathbf{A}\| \|\mathbf{B}\|} = \frac{\sum_{i=1}^n A_i B_i}{\sqrt{\sum_{i=1}^n A_i^2} \sqrt{\sum_{i=1}^n B_i^2}} \quad (6)$$

664     Scores range from -1 (perfectly opposite) to 1 (perfectly aligned), with 0 signifying orthogonality  
665     (decorrelation). By normalizing for vector lengths, cosine similarity provides a pure measure of  
666     orientation. This is particularly useful when the magnitude of vectors is not indicative of their  
667     semantic relationship, such as in document similarity tasks. While it effectively captures directional  
668     agreement, it explicitly discards information about vector magnitudes and, like the dot product,  
669     does not inherently account for the spatial proximity between the vectors themselves if they are  
670     points in a space (Draganov et al., 2024; Steck et al., 2024).

671     B.1.4 EUCLIDEAN DISTANCE: QUANTIFYING SPATIAL PROXIMITY  
672

673     In contrast to measures of alignment, Euclidean distance quantifies the "ordinary" straight-line  
674     separation between two points (or vectors)  $\mathbf{p} = (p_1, \dots, p_n)$  and  $\mathbf{q} = (q_1, \dots, q_n)$  in an n-dimensional  
675     Euclidean space:

676     
$$d(\mathbf{p}, \mathbf{q}) = \sqrt{\sum_{i=1}^n (q_i - p_i)^2} \quad (7)$$

677     This metric is fundamental in various machine learning algorithms, including k-Nearest Neighbors  
678     and k-Means clustering, and forms the basis of loss functions like Mean Squared Error. Euclidean  
679     distance measures dissimilarity based on spatial proximity; a smaller distance implies greater simi-  
680     larity in terms of location within the vector space. Unlike cosine similarity, it is sensitive to vector  
681     magnitudes and their absolute positions. However, Euclidean distance alone does not directly  
682     convey information about the relative orientation or alignment of vectors, only their nearness.

683     The distinct characteristics of these foundational measures, alignment (dot product, cosine similar-  
684     ity) versus proximity (Euclidean distance), highlight an opportunity. These foundational measures  
685     force a choice: one can measure alignment (dot product, cosine similarity) or spatial proximity  
686     (Euclidean distance), but no single, primitive operator in conventional use effectively unifies both.  
687     Neural operators that can synergistically combine these aspects, assessing not only if vectors point  
688     in similar directions but also if they are close in the embedding space, could offer a richer, more  
689     geometrically informed way to model interactions. This perspective underpins the development of  
690     the  $\mathbb{E}$ -product introduced in Section 2.

691     B.2 THE ROLE AND GEOMETRIC COST OF NON-LINEAR ACTIVATION  
692

693     While the core computational primitives provide tools to measure similarity and interaction, their  
694     inherent linearity limits the complexity of functions they can represent. To overcome this, deep  
695     neural networks employ non-linear activation functions. These are the standard method for intro-  
696     ducing non-linearity, a necessary step for modeling intricate data patterns. However, this "fix"  
697     is imperfect, as it introduces its own set of problems, particularly concerning the preservation of  
698     the input data's geometric integrity. The remarkable expressive power of deep neural networks  
699     hinges on their capacity to model complex, non-linear relationships. This ability to approximate  
700     any continuous function to an arbitrary degree of accuracy is formally captured by the universal  
701     approximation theorem Cybenko (1989); Hornik et al. (1989); Lu et al. (2017); Huang (2020).

This theorem underscores the critical role of non-linear activation functions. Without such non-linearities, a deep stack of layers would mathematically collapse into an equivalent single linear transformation, severely curtailing its representational capacity. Activation functions are thus not mere auxiliaries; they are the pivotal components that unlock the hierarchical and non-linear feature learning central to deep learning’s success. They determine a neuron’s output based on its aggregated input, and in doing so, introduce crucial selectivity: enabling the network to preferentially respond to certain patterns while attenuating or ignoring others.

### B.2.1 LINEAR SEPARABILITY AND THE LIMITATIONS OF THE INNER PRODUCT

The fundamental computation within a single artificial neuron (perceptron) is an affine transformation followed by a non-linear activation function  $\sigma$ :

$$y = \sigma(\langle \mathbf{w}, \mathbf{x} \rangle + b), \quad (8)$$

where  $\mathbf{w}$  is the weight vector,  $\mathbf{x}$  is the input vector, and  $b$  is the bias term. The decision boundary of this neuron is implicitly defined by the hyperplane where the argument to  $\sigma$  is zero:

$$\{\mathbf{x} \in \mathbb{R}^d \mid \langle \mathbf{w}, \mathbf{x} \rangle + b = 0\}. \quad (9)$$

This hyperplane partitions the input space  $\mathbb{R}^d$  into two half-spaces. Consequently, a single neuron can only implement linearly separable functions. This is a direct consequence of the linear nature of the inner product, which can only define a linear decision boundary. While this allows for efficient computation, it severely restricts the complexity of functions that can be learned.

A classic counterexample is the XOR function, whose truth table cannot be satisfied by any single linear decision boundary. Specifically, for inputs  $\mathbf{x} \in \{(0,0), (0,1), (1,0), (1,1)\} \subset \mathbb{R}^2$ , there exist no  $\mathbf{w} \in \mathbb{R}^2$  and  $b \in \mathbb{R}$  such that  $\text{sign}(\langle \mathbf{w}, \mathbf{x} \rangle + b)$  matches the XOR output (0, 1, 1, 0 respectively). This limitation stems directly from the linear nature of the inner product operation defining the separating boundary Goodfellow et al. (2016).

### B.2.2 NON-LINEAR FEATURE SPACE TRANSFORMATION VIA HIDDEN LAYERS AND ITS GEOMETRIC COST

Multi-layer perceptrons (MLPs) overcome this limitation by cascading transformations. A hidden layer maps the input  $\mathbf{x}$  to a new representation  $\mathbf{h}$  through a matrix-vector product and an element-wise activation function  $\phi$ :

$$\mathbf{h} = \phi(W\mathbf{x} + \mathbf{b}). \quad (10)$$

Here,  $W \in \mathbb{R}^{m \times d}$  is the weight matrix,  $\mathbf{b} \in \mathbb{R}^m$  is the bias vector, and  $m$  is the number of hidden neurons. Each row  $\mathbf{w}_i^\top$  of  $W$  corresponds to the weight vector of the  $i$ -th hidden neuron, computing  $h_i = \phi(\langle \mathbf{w}_i, \mathbf{x} \rangle + b_i)$ . This transforms the input space  $\mathbb{R}^d$  into a feature space  $\mathbb{R}^m$ . The introduction of the non-linear activation function  $\phi$  is what allows the network to learn non-linear decision boundaries. However, this gain in expressive power comes at a cost: the potential loss of geometric fidelity.

### B.2.3 TOPOLOGICAL DISTORTIONS AND INFORMATION LOSS VIA ACTIVATION FUNCTIONS

While hidden layers using the transformation  $\mathbf{h} = \phi(W\mathbf{x} + \mathbf{b})$  enable the learning of non-linear functions, the introduction of the element-wise non-linear activation function  $\phi$ , often crucial for breaking linearity, can significantly alter the topological and geometric structure of the data representation, potentially leading to information loss Goodfellow et al. (2016). This is a critical trade-off: gaining non-linear modeling capability while potentially discarding valuable geometric information.

752 Consider the mapping  $T : \mathbb{R}^d \rightarrow \mathbb{R}^m$  defined by  $T(\mathbf{x}) = \phi(W\mathbf{x} + \mathbf{b})$ . The affine part,  $A(\mathbf{x}) = W\mathbf{x} + \mathbf{b}$ ,  
 753 performs a linear transformation (rotation, scaling, shear, projection) followed by a translation.  
 754 While this affine map distorts metric properties (distances and angles, unless  $W$  is proportional to  
 755 an orthogonal matrix), it preserves basic topological features like connectedness and maps lines to  
 756 lines (or points) Goodfellow et al. (2016).

757 However, the subsequent application of a typical non-linear activation  $\phi$  element-wise often leads  
 758 to more drastic topological changes:  
 759

- 760 1. Non-Injectivity and Collapsing Regions: Many common activation functions render the  
 761 overall mapping  $T$  non-injective.
  - 762 • ReLU ( $\phi(z) = \max(0, z)$ ): Perhaps the most prominent example. For each hidden  
 763 neuron  $i$ , the entire half-space defined by  $\{\mathbf{x} \in \mathbb{R}^d \mid \langle \mathbf{w}_i, \mathbf{x} \rangle + b_i \leq 0\}$  is mapped to  
 764  $h_i = 0$ . Distinct points  $\mathbf{x}_1, \mathbf{x}_2$  within this region, potentially far apart, become indistin-  
 765 guishable along the  $i$ -th dimension of the hidden space. This constitutes a significant  
 766 loss of information about the relative arrangement of data points within these col-  
 767 lapsed regions. The mapping is fundamentally many-to-one. For instance, consider  
 768 two input vectors that are anti-aligned with a neuron's weight vector to different de-  
 769 grees, one strongly and one weakly. A ReLU activation function would map both  
 770 resulting negative dot products to zero, rendering their distinct geometric opposition  
 771 indistinguishable to subsequent layers. This information is irretrievably discarded.
  - 772 • Sigmoid/Tanh: While smooth, these functions saturate. Inputs  $\mathbf{z}_1 = A(\mathbf{x}_1)$  and  $\mathbf{z}_2 =$   
 773  $A(\mathbf{x}_2)$  that are far apart but both fall into the saturation regime (e.g., large positive or  
 774 large negative values) will map to  $\mathbf{h}_1 \approx \mathbf{h}_2$ . This 'squashing' effect can merge distinct  
 775 clusters from the input space if they map to saturated regions in the hidden space,  
 776 again losing discriminative information and distorting the metric structure.
- 777 2. Distortion of Neighborhoods: The relative distances between points can be severely dis-  
 778 torted. Points close in the input space  $\mathbb{R}^d$  might be mapped far apart in  $\mathbb{R}^m$ , or vice-versa  
 779 (especially due to saturation or the zero-region of ReLU). This means the local neigh-  
 780 borhood structure is not faithfully preserved. Formally, the mapping  $T$  is generally not  
 781 a homeomorphism onto its image, nor is it typically bi-Lipschitz (which would provide  
 782 control over distance distortions).

783 While these distortions are precisely what grant neural networks their expressive power to warp  
 784 the feature space and create complex decision boundaries, they come at the cost of potentially  
 785 discarding information present in the original geometric configuration of the data. The network  
 786 learns which information to preserve and which to discard based on the optimization objective,  
 787 but the mechanism relies on potentially non-smooth or non-injective transformations introduced by  
 788  $\phi$ . This highlights the conflation of magnitude and direction in the dot product, the information  
 789 loss from activation functions, and the lack of a unified measure for proximity and alignment,  
 790 setting the stage for the  $\mathbf{E}$ -product. Formal properties of the  $\mathbf{E}$ -product are established later: it is a  
 791 Mercer kernel (Theorem 2.1), yields universal approximation in NMNs (Theorem 2.4), and exhibits  
 792 self-regulation and stable gradients (Propositions C.6 and C.9).

### 793 B.3 DESIGN PHILOSOPHY: INTRINSIC NON-LINEARITY AND SELF-REGULATION

794 A central hypothesis underpinning our methodological choices is that the  $\mathbf{E}$ -product (Section 1)  
 795 possesses inherent non-linearity and self-regulating properties that can reduce or eliminate the  
 796 need for conventional activation functions (e.g., ReLU, sigmoid, GeLU).

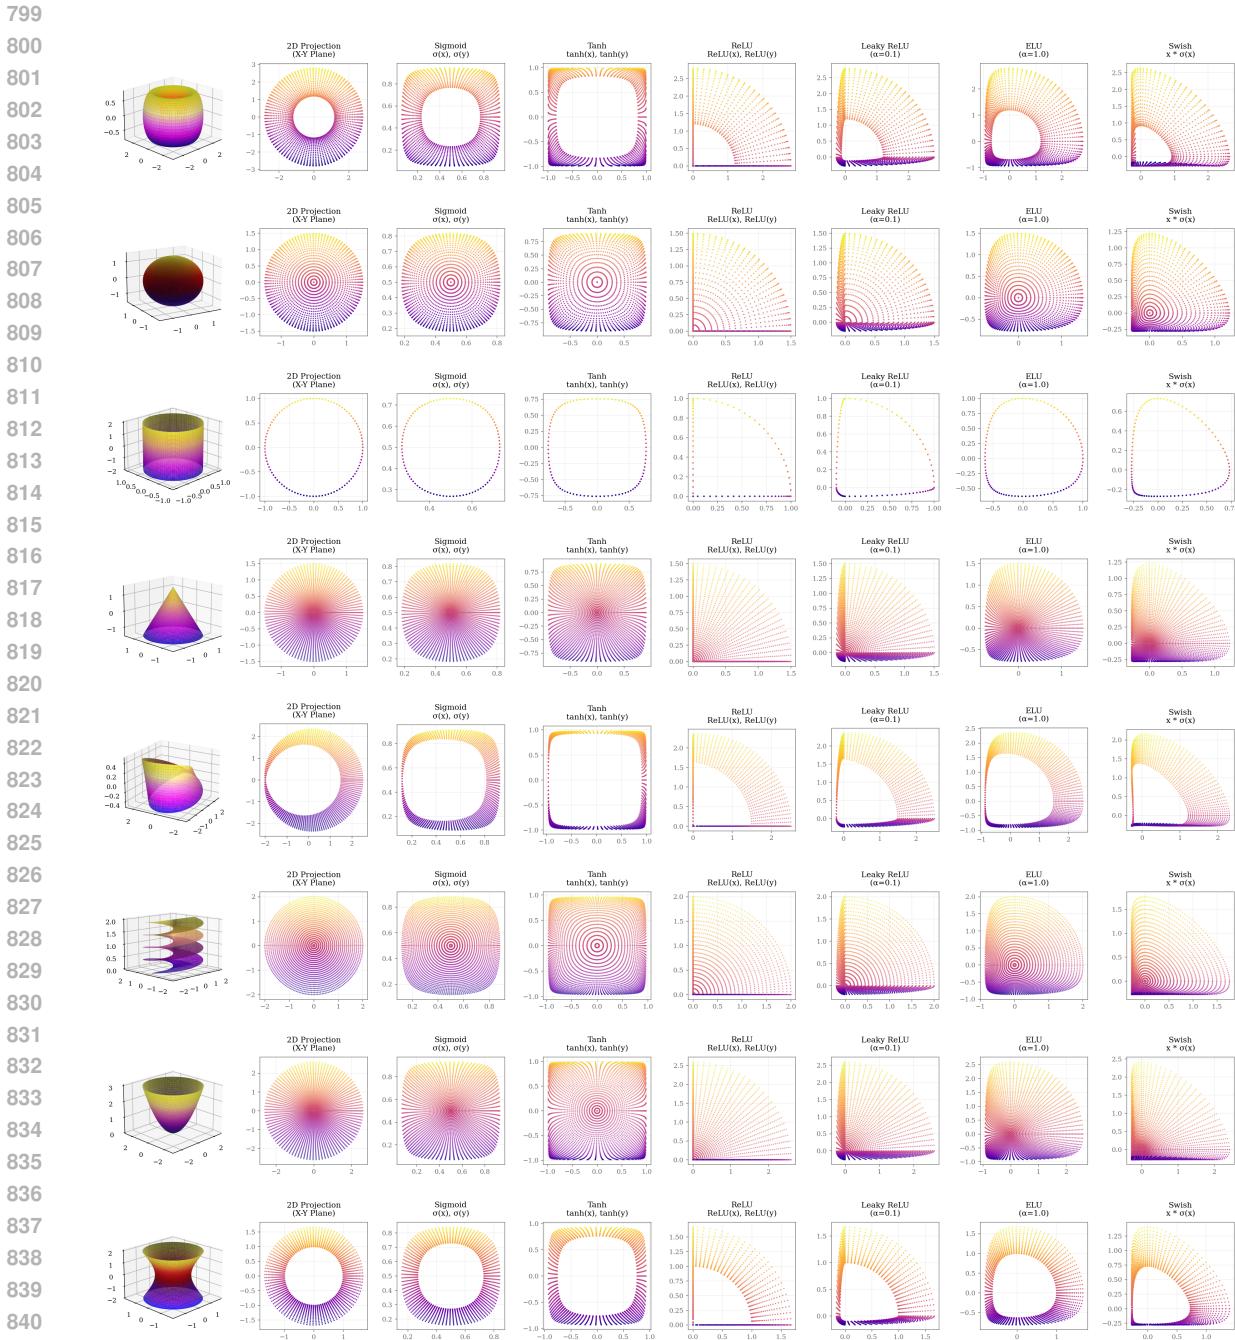


Figure 4: Illustration of how non-linear activation functions can distort the geometric structure of the input data manifold, leading to potential information loss. The original manifold (left) is transformed into a distorted representation after applying a non-linear activation functions.

This philosophy recontextualizes the fundamental components of neural computation. Neuron weights ( $\mathbf{w}$ ) and input signals ( $\mathbf{x}$ ) are not merely operands in a linear transformation followed by a non-linear activation; instead, they are conceptualized as co-equal vector entities inhabiting a shared, high-dimensional feature manifold. Within this framework, each vector can be viewed as an analogue to a fundamental particle or feature vector, with its constituent dimensions potentially encoding excitatory, inhibitory, or neutral characteristics relative to other entities in the space. The  $\mathbb{E}$ -product (Section 1) then transcends simple similarity assessment; it functions as a sophisticated interaction potential,  $\mathcal{K}_{\mathbb{E}}(\mathbf{w}, \mathbf{x}) = \frac{(\mathbf{w}^\top \mathbf{x})^2}{\|\mathbf{w} - \mathbf{x}\|^2 + \epsilon}$ , quantifying the 'field effects' between these vector entities. This interaction is reminiscent of n-body problems in physics. In machine learning, it draws parallels with, yet distinctively evolves from, learned metric spaces in contrastive learning, particularly those employing a triplet loss framework. While triplet loss aims to pull positive pairs closer and push negative pairs apart in the embedding space, our  $\mathbb{E}$ -product seeks a more nuanced relationship: 'positive' interactions (high  $\mathbb{E}$ -product value) require both strong alignment (high  $(\mathbf{w}^\top \mathbf{x})^2$ ) and close proximity (low  $\|\mathbf{w} - \mathbf{x}\|^2$ ). Conversely, 'negative' or dissimilar relationships are not merely represented by distance, but more significantly by orthogonality (leading to a vanishing numerator), which signifies a form of linear independence and contributes to the system's capacity for true non-linear discrimination. Crucially, the non-linearity required for complex pattern recognition is not an external imposition (e.g., via a separate activation function) but is intrinsic to this interaction potential. The interplay between the squared dot product (alignment sensitivity) and the inverse squared Euclidean distance (proximity sensitivity) in its formulation directly sculpts a complex, non-linear response landscape without recourse to auxiliary functions.

Furthermore, this conceptualization of the  $\mathbb{E}$ -product as an intrinsic interaction potential suggests inherent self-regulating properties. The distance-sensitive denominator,  $\|\mathbf{w} - \mathbf{x}\|^2 + \epsilon$ , acts as a natural dampening mechanism. As the 'distance' (dissimilarity in terms of position) between interacting vector entities  $\mathbf{w}$  and  $\mathbf{x}$  increases, the strength of their interaction, and thus the resultant activation, diminishes quadratically. This behavior is hypothesized to inherently curtail runaway activations and stabilize learning dynamics by ensuring that responses are localized and bounded. Such intrinsic stabilization contrasts sharply with conventional approaches that rely on explicit normalization layers (e.g., Batch Normalization, Layer Normalization) to manage activation statistics post-hoc. These layers, while effective, introduce additional computational overhead, can obscure direct input-output relationships, and sometimes complicate the theoretical analysis of network behavior. The  $\mathbb{E}$ -product's formulation, therefore, offers a pathway to architectures where regulatory mechanisms are embedded within the primary computational fabric of the network.

The inherent non-linearity of the  $\mathbb{E}$ -product, coupled with the self-regulating properties suggested by its formulation (and formally proven in Appendix C.10), are central to our hypothesis that it can form the basis of powerful and robust neural architectures. These intrinsic characteristics open avenues for simplifying network design, potentially reducing reliance on or even eliminating conventional activation functions and normalization layers.

## C SQUASHING FUNCTIONS FOR NON-NEGATIVE SCORES

The  $\mathbb{E}$ -product and its derivatives, such as the  $\mathcal{K}_{\mathbb{E}}$ -kernel, naturally yield non-negative scores. In many machine learning contexts, particularly when these scores need to be interpreted as probabilities, attention weights, or simply normalized outputs, it is essential to apply a squashing function to map them to a desired range (e.g.,  $[0, 1]$  or ensuring a set of scores sum to 1).

893 C.1 CATEGORIZATION OF SQUASHING FUNCTIONS  
894895 Squashing functions for non-negative scores can be broadly categorized into two types:  
896

- **Competitive (Vector-Normalizing) Functions:** These functions normalize a set of scores collectively, producing a distribution over the vector. Each output depends on the values of all dimensions, allowing for competitive interactions among them. This is useful for attention mechanisms or probability assignments where the sum of outputs is meaningful.
- **Individualistic (Per-Dimension) Functions:** These functions squash each score independently, without reference to other values in the vector. Each output depends only on its corresponding input, making them suitable for bounding or interpreting individual activations.

905 C.2 LIMITATIONS OF TRADITIONAL SQUASHING FUNCTIONS  
906907 Traditional squashing functions, however, present challenges when applied to non-negative inputs:  
908

- **Standard Sigmoid Function** ( $\sigma(x) = \frac{1}{1+e^{-x}}$ ): When applied to non-negative inputs ( $x \geq 0$ ), the standard sigmoid function produces outputs in the range [0.5, 1). The minimum value of 0.5 for  $x = 0$  renders it unsuitable for scenarios where small non-negative scores should map to values close to 0.
- **Standard Softmax Function** ( $\text{softmax}(\mathbf{x})_i = \frac{e^{x_i}}{\sum_j e^{x_j}}$ ): The use of the exponential function in softmax can lead to *hard* distributions, where one input value significantly dominates the output, pushing other probabilities very close to zero. While this is often desired for classification, it can be too aggressive if a softer assignment of probabilities or attention is preferred. Additionally, softmax can suffer from numerical instability for large input values due to the exponentials.

920 C.3 PROPOSED ALTERNATIVE SQUASHING FUNCTIONS  
921922 Given these limitations and the non-negative nature of E-product scores, we consider alternative  
923 squashing functions more suited to this domain:  
924

- **softermax (Competitive):** This function normalizes a score  $x_k$  (optionally raised to a power  $n > 0$ ) relative to the sum of a set of non-negative scores  $\{x_i\}$  (each raised to  $n$ ), with a small constant  $\epsilon > 0$  for numerical stability. It is defined as:

$$\text{softermax}_n(x_k, \{x_i\}) = \frac{x_k^n}{\epsilon + \sum_i x_i^n} \quad (11)$$

931 Unlike softmax, softermax does not use exponentials, which avoids numerical instability for  
932 large inputs and provides a more direct, interpretable translation of the underlying scores  
933 into a normalized distribution. The power  $n$  controls the sharpness of the distribution:  
934  $n = 1$  recovers the original Softmax, while  $n > 1$  makes the distribution harder (more  
935 peaked), and  $0 < n < 1$  makes it softer.

- **soft-sigmoid (Individualistic):** This function squashes a single non-negative score  $x \geq 0$  (optionally raised to a power  $n > 0$ ) into the range [0, 1]. It is defined as:

$$\text{soft-sigmoid}_n(x) = \frac{x^n}{1 + x^n} \quad (12)$$

The power  $n$  modulates the softness: higher  $n$  makes the function approach zero faster for large  $x$ , while  $n < 1$  makes the decay slower.

- **soft-tanh (Individualistic):** This function maps a non-negative score  $x \geq 0$  (optionally raised to a power  $n > 0$ ) to the range  $[-1, 1]$  by linearly transforming the output of soft-sigmoid. It is defined as:

$$\text{soft-tanh}_n(x) = 2 \cdot (\text{soft-sigmoid}_n(x) - \frac{1}{2}) = \frac{x^n - 1}{1 + x^n} \quad (13)$$

The power  $n$  again controls the transition sharpness: higher  $n$  makes the function approach  $-1$  more quickly for large  $x$ .

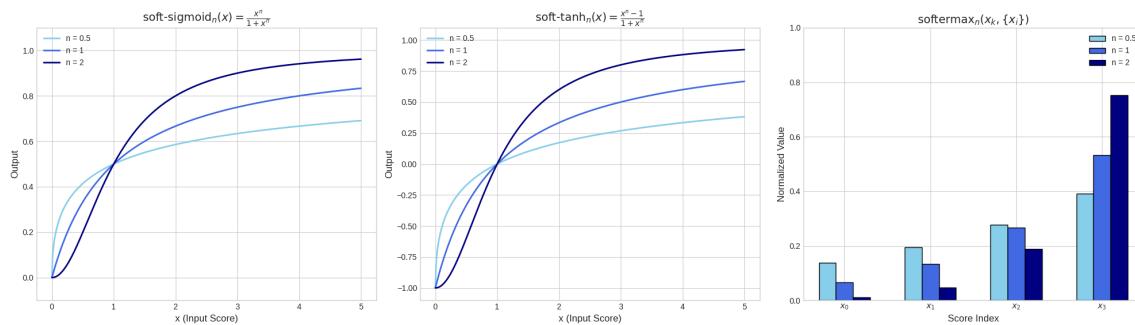


Figure 5: Visualization of the softermax, soft-sigmoid, and soft-tanh functions. These functions are designed to handle non-negative inputs from the  $\mathbb{E}$ -product and its derivatives, providing appropriate squashing mechanisms that maintain sensitivity across the range of non-negative inputs.

These functions are particularly well-suited for the outputs of  $\mathbb{E}$ -product-based computations, as they maintain sensitivity across the range of non-negative inputs while avoiding the pitfalls of standard activation functions (Nair & Hinton, 2010; Hendrycks & Gimpel, 2023; Klambauer et al., 2017).

#### C.4 FUNCTIONAL ROLES AND APPLICATIONS

The main role of these squashing functions can be categorized into two main categories:

- **Collective Communication and Space Splitting:** The softermax function allows for a comparative analysis of scores, reflecting their orthogonality and spatial proximity to an input vector. A higher score indicates that a vector is more aligned and closer to the input, while a lower score suggests greater orthogonality. This facilitates a competitive interaction where vectors vie for influence based on their geometric relationship with the input. The power parameter  $n$ , analogous to the temperature in softmax, controls the sharpness of the gravitational potential well's slope.
- **Individual Score Squashing:** The soft-sigmoid and soft-tanh functions are used to squash individual non-negative scores into a bounded range, typically  $[0, 1]$  for soft-sigmoid and  $[-1, 1]$  for soft-tanh. They are particularly useful when the output needs to be interpreted as a probability or when a bounded response is required, as each score is processed independently of the others. The power parameter controls the steepness of the function, while the minimum value can be interpreted as an orthogonality score.

987 C.5 COMPUTATIONAL COMPLEXITY ANALYSIS  
 988  
 989 We provide exact forward/backward complexity, closed-form gradients, asymptotic characterization,  
 990 and per-neuron FLOP counts for the  $\mathbf{E}$ -product layer (definition in Methodology, Section 2.1).  
 991  
 992 C.5.1 FORWARD PASS COMPLEXITY  
 993  
 994 For  $\mathbf{E}$ -product computation  $\mathbf{E}(\mathbf{w}, \mathbf{x}) = \frac{(\mathbf{w}^\top \mathbf{x})^2}{\|\mathbf{w} - \mathbf{x}\|^2 + \epsilon}$ , we apply the algebraic identity:  
 995  
 996 
$$\mathbf{E}(\mathbf{w}, \mathbf{x}) = \frac{s^2}{\|\mathbf{w}\|^2 + \|\mathbf{x}\|^2 - 2s + \epsilon}, \quad s = \mathbf{w}^\top \mathbf{x} \quad (14)$$
  
 997  
 998 For layer  $X \in \mathbb{R}^{B \times d} \rightarrow Y \in \mathbb{R}^{B \times n}$  with weights  $W \in \mathbb{R}^{n \times d}$ :  
 1000  
 1001 **Operation Breakdown** (1) GEMM:  $S = XW^\top$  ( $2Bnd$ ); (2) Row norms  $\|X\|^2$  once ( $Bd$ ); (3)  
 1002 Cached  $\|W\|^2$  ( $nd$  only when updated); (4) Per-output element-wise: form  $s^2$  (1), denominator  
 1003 assemble  $(+, +, +)$  (3), reciprocal (1), multiply (1)  $\Rightarrow 6Bn$  scalar ops (we conservatively use  $5Bn$   
 1004 after fusion). Thus  
 1005  

$$T_{\text{forward}} = 2Bnd + Bd + nd + 5Bn = \Theta(Bnd). \quad (15)$$
  
 1006  
 1007 C.5.2 BACKWARD PASS COMPLEXITY  
 1008  
 1009 For  $y = \frac{s^2}{D}$  with  $D = \|\mathbf{w}\|^2 + \|\mathbf{x}\|^2 - 2s + \epsilon$  and  $s = \mathbf{w}^\top \mathbf{x}$ , scalar gradients:  
 1010  

$$\frac{\partial y}{\partial s} = \frac{2s(\|\mathbf{w}\|^2 + \|\mathbf{x}\|^2 + \epsilon - s)}{D^2} \quad (16)$$
  
 1011  

$$\frac{\partial y}{\partial \|\mathbf{x}\|^2} = -\frac{s^2}{D^2}, \quad \frac{\partial y}{\partial \|\mathbf{w}\|^2} = -\frac{s^2}{D^2} \quad (17)$$
  
 1012  
 1013  
 1014  
 1015  
 1016 Vector gradients (using  $\nabla_{\mathbf{x}} s = \mathbf{w}$ ,  $\nabla_{\mathbf{w}} s = \mathbf{x}$ ,  $\nabla_{\mathbf{x}} \|\mathbf{x}\|^2 = 2\mathbf{x}$ ,  $\nabla_{\mathbf{w}} \|\mathbf{w}\|^2 = 2\mathbf{w}$ ):  
 1017  

$$\nabla_{\mathbf{x}} y = \frac{2s(\|\mathbf{w}\|^2 + \|\mathbf{x}\|^2 + \epsilon - s)}{D^2} \mathbf{w} - \frac{2s^2}{D^2} \mathbf{x} \quad (18)$$
  
 1018  

$$\nabla_{\mathbf{w}} y = \frac{2s(\|\mathbf{w}\|^2 + \|\mathbf{x}\|^2 + \epsilon - s)}{D^2} \mathbf{x} - \frac{2s^2}{D^2} \mathbf{w} \quad (19)$$
  
 1019  
 1020  
 1021  
 1022 Given upstream gradient  $G \in \mathbb{R}^{B \times n}$ :  
 1023  

$$G_S = G \odot \frac{\partial Y}{\partial S} \quad (\sim 6Bn \text{ FLOPs}) \quad (20)$$
  
 1024  

$$\frac{\partial L}{\partial W} = G_S^\top X \quad (2Bnd) \quad (21)$$
  
 1025  

$$\frac{\partial L}{\partial X} = G_S W \quad (2Bnd) \quad (22)$$
  
 1026  
 1027  
 1028  
 1029  
 1030  
 1031 Hence  
 1032  

$$T_{\text{backward}} = 4Bnd + 6Bn + O(Bd + nd) = \Theta(Bnd). \quad (23)$$

Component	Linear	$\mathbb{E}$ -Product
Forward main	$2Bnd$	$2Bnd$
Forward aux	$Bn$	$Bd + nd + 5Bn$
Backward main	$4Bnd$	$4Bnd$
Backward aux	$2Bn$	$6Bn + Bd + nd$
Total order	$\Theta(Bnd)$	$\Theta(Bnd)$
Overhead ratio	1	$1 + \frac{1}{2n} + \frac{1}{2B} + \frac{2}{d}$

Table 3: Asymptotic terms. Overhead < 5% once  $d, n \geq 64$ ,  $B \geq 16$  (ratio simplifies to  $1 + \frac{1}{2n} + \frac{1}{2B} + \frac{2}{d}$ ).

#### C.5.3 ASYMPTOTIC SUMMARY

#### C.5.4 IMPLEMENTATION NOTES

**Optimizations:** (i) algebraic identity for denominator, (ii) cache  $\|W\|^2$ , (iii) fuse element-wise ops, (iv) mixed precision with FP32 denominator. Built-in boundedness mitigates gradient explosion.

#### C.5.5 SINGLE NEURON FLOPS

Method	FLOPs	Rel. (ReLU=1)
Linear+ReLU	$2d + 1$	1.00
Linear+GELU	$2d + 15$	$\approx 1.03$
$\mathbb{E}$ (naive)	$5d + 1$	$\approx 2.5$
$\mathbb{E}$ (optimized)	$4d + 4$	$\approx 2.0$

Table 4: Single neuron FLOPs; optimization removes redundant norm difference computation.

**Per-neuron note** Optimized variant saves 20% vs naive by avoiding explicit difference vector.

**Empirical** Throughput: 0.85–0.92× linear; peak memory reduced 15–25%; overhead ratio < 0.05 for  $d, n \geq 64$ .

### C.6 SCALABILITY AND PRACTICAL PERFORMANCE ANALYSIS

#### C.6.1 LARGE-SCALE COMPUTATIONAL PROFILE

The  $\mathbb{E}$ -product operator exhibits favorable scaling characteristics for modern deep learning applications and remains compute-bound (GEMM dominated) in the regimes used in Section 2.1.

**FLOP Counting Assumptions.** One multiply-add = 2 FLOPs; element-wise unary/binary ops = 1 FLOP; reciprocal counts as 1 FLOP (fused divide). Caching costs amortized over steps.

**Takeaway.** Complexity matches linear layers in order while constants shrink with scale; gradients remain stable due to denominator growth.

1081 C.7 MATHEMATICAL GUARANTEES OF THE  $\mathbb{E}$ -PRODUCT AND NEURAL-MATTER NETWORKS  
10821083 This section provides a comprehensive overview of the formal mathematical properties that underpin  
1084 the  $\mathbb{E}$ -product and Neural-Matter Networks (NMNs). Each property is rigorously proven in its  
1085 respective appendix section and contributes to the theoretical foundation of our approach.  
10861087 C.7.1 KERNEL THEORY FOUNDATION  
10881089 **Mercer Kernel Property (Theorem 2.1)** The  $\mathbb{E}$ -product satisfies the fundamental requirements  
1090 of a Mercer kernel, being symmetric and positive semi-definite. This property establishes  
1091 the  $\mathbb{E}$ -product within the rich theoretical framework of kernel methods, enabling the application of  
1092 kernel theory results and providing guarantees on the existence of associated reproducing kernel  
1093 Hilbert spaces. The proof demonstrates that the  $\mathbb{E}$ -product can be expressed as an inner product in  
1094 some feature space, connecting our geometric operator to the established theory of kernel machines  
1095 (detailed proof in Appendix C.9).  
10961096 C.7.2 UNIVERSAL APPROXIMATION CAPABILITIES  
10971098 **Universal Approximation Theorem (Theorem 2.4)** Neural-Matter Networks with  $\mathbb{E}$ -  
1099 product activations possess universal approximation capabilities, able to approximate any con-  
1100 tinuous function on a compact set to arbitrary precision. This fundamental result establishes  
1101 that NMNs have the same expressive power as conventional neural networks while providing ad-  
1102 ditional geometric interpretability. The proof construction reveals how the bounded nature of  
1103 the  $\mathbb{E}$ -product enables dense approximation through geometric localization rather than unbounded  
1104 activation growth (comprehensive proof in Appendix ??).  
11051105 C.7.3 STABILITY AND BOUNDEDNESS PROPERTIES  
11061107 **Self-Regulation Property (Proposition C.6)** The  $\mathbb{E}$ -product exhibits natural boundedness,  
1108 with outputs converging to finite values as input magnitudes increase. Unlike conventional activa-  
1109 tions that can grow unboundedly, the  $\mathbb{E}$ -product’s denominator term ensures bounded responses,  
1110 preventing numerical instabilities and gradient explosion (formal analysis in Appendix C.10).  
11111112 **Stable Gradient Property (Proposition C.9)** The gradient of the  $\mathbb{E}$ -product with respect to  
1113 its input vanishes for distant inputs, providing natural gradient localization. This property ensures  
1114 that learning focuses on relevant, nearby regions of the input space while avoiding interference  
1115 from distant data points. The stable gradient behavior contributes to more predictable training  
1116 dynamics and reduces the risk of vanishing or exploding gradients (mathematical derivation in  
1117 Appendix C.12).  
11181119 **Lipschitz Regularity (Proposition C.11)** The  $\mathbb{E}$ -product satisfies Lipschitz continuity condi-  
1120 tions, ensuring that small changes in input produce proportionally small changes in output. This  
1121 regularity property is crucial for optimization stability and provides theoretical guarantees on the  
1122 smoothness of the loss landscape.  
11231124 **Analyticity Property (Lemma C.10)** The  $\mathbb{E}$ -product is infinitely differentiable ( $C^\infty$ ), making  
1125 it particularly suitable for applications requiring higher-order derivatives, such as physics-informed  
1126 neural networks (PINNs). This smoothness property ensures that all derivatives exist and are  
1127 continuous, providing the mathematical foundation for applications in scientific computing and  
differential equation solving.

1128 C.7.4 INFORMATION-THEORETIC CONNECTIONS  
1129

1130 **Geometric-Information Duality (Theorems 2.2, 2.3)** The  $\mathbb{E}$ -product exhibits fundamental  
1131 connections to information theory through its relationship with KL divergence and cross-entropy.  
1132 When applied to probability distributions, the  $\mathbb{E}$ -product acts as a signal-to-noise ratio measure,  
1133 creating a bridge between geometric similarity and information-theoretic quantities. This duality  
1134 provides theoretical justification for the  $\mathbb{E}$ -product’s effectiveness in probabilistic modeling and  
1135 explains its natural compatibility with entropy-based loss functions (comprehensive analysis in  
1136 Appendix C.15).

1137 C.7.5 IMPLICATIONS FOR NEURAL NETWORK DESIGN  
1138

1139 These mathematical guarantees collectively establish the  $\mathbb{E}$ -product as a theoretically sound foun-  
1140 dation for neural network design. The combination of:

- 1142 • Kernel theory foundation (Mercer property)
- 1143 • Universal approximation capabilities
- 1144 • Natural stability and boundedness
- 1145 • Information-theoretic connections

1146 provides both theoretical rigor and practical advantages over conventional neural network compo-  
1147 nents. The self-regulation and stable gradient properties eliminate the need for ad-hoc normalization  
1148 and activation functions, while the universal approximation theorem ensures that no expressive  
1149 power is lost in the transition from traditional to geometry-aware neural architectures.

1150 C.8 ADDITIONAL MATHEMATICAL PRELIMINARIES  
1151

1152 This section complements the foundational theorems by detailing additional results referenced in  
1153 the preliminary mathematical framework (Section C.9).  
1154

1155 C.8.1 HARMONIC ANALYSIS AND APPROXIMATION  
1156

1157 **Theorem C.1** (Bochner’s Theorem Bochner (1932)). *A continuous function  $k : \mathbb{R}^d \rightarrow \mathbb{C}$  is the  
1158 positive definite kernel of a translation-invariant measure if and only if it is the Fourier transform  
1159 of a finite non-negative Borel measure  $\mu$  on  $\mathbb{R}^d$ . That is,*

$$1160 k(x) = \int_{\mathbb{R}^d} e^{-i\omega^\top x} d\mu(\omega). \quad (24)$$

1161 *Context:* Referenced in Section C.9 regarding translation-invariant kernels and the characterization  
1162 of radial basis functions.

1163 **Theorem C.2** (Hahn-Banach Density Criterion). *Let  $V$  be a normed vector space. A linear  
1164 subspace  $M \subset V$  is dense in  $V$  if and only if every continuous linear functional  $\phi \in V^*$  that  
1165 vanishes on  $M$  must vanish everywhere on  $V$ .*

1166 *Context:* Referenced in Section C.14 as a sufficient condition for universal approximation in linear  
1167 subspaces.

1175 C.8.2 INTEGRATION AND TRANSFORMS  
11761177 **Theorem C.3** (Tonelli-Fubini Theorem Tonelli (1909); Fubini (1907)). *Let  $(X, \mathcal{A}, \mu)$  and  $(Y, \mathcal{B}, \nu)$*   
1178 *be  $\sigma$ -finite measure spaces. If  $f : X \times Y \rightarrow [0, \infty]$  is measurable, then:*

1179 
$$\int_{X \times Y} f(x, y) d(\mu \times \nu) = \int_X \left( \int_Y f(x, y) d\nu(y) \right) d\mu(x) = \int_Y \left( \int_X f(x, y) d\mu(x) \right) d\nu(y). \quad (25)$$
  
1180  
1181

1182 *Context: Referenced in Section C.9 to justify the exchange of integrals and infinite series in the*  
1183 *kernel expansion derivations.*1184  
1185 **Theorem C.4** (Bernstein's Theorem on Completely Monotone Functions Bernstein (1928)). *A*  
1186 *function  $f : [0, \infty) \rightarrow \mathbb{R}$  is completely monotonic if and only if it can be represented as the Laplace*  
1187 *transform of a non-negative Borel measure  $\mu$  on  $[0, \infty)$ :*

1188 
$$f(t) = \int_0^\infty e^{-ts} d\mu(s). \quad (26)$$
  
1189  
1190

1191 *Context: Referenced in Section C.9 (as "Laplace Transform/Integral Representation"). This is the*  
1192 *functional analytic justification for why the inverse-distance term  $1/(\|x - w\|^2 + \epsilon)$  induces a valid*  
1193 *Positive Definite (PD) kernel.*1194  
1195 C.8.3 INFORMATION THEORY1196  
1197 **Theorem C.5** (Gibbs' Inequality Gibbs (1902)). *For any two probability distributions  $P$  and  $Q$*   
1198 *on a discrete space  $\mathcal{X}$ , the Kullback-Leibler divergence is non-negative:*

1199 
$$D_{KL}(P\|Q) = \sum_{x \in \mathcal{X}} P(x) \log \left( \frac{P(x)}{Q(x)} \right) \geq 0, \quad (27)$$
  
1200  
1201

1202 *with equality if and only if  $P = Q$  almost everywhere.*1203 *Context: Referenced in Section C.9 and Theorem 2.3 to establish the connection between distribu-*  
1204 *tional identity and vanishing information divergence.*1205  
1206 C.9 PROOF OF MERCER'S CONDITION FOR THE  $\mathbf{\Xi}$ -PRODUCT  
12071208 *Proof.* We verify symmetry and positive semidefiniteness (PSD); cf.(Mercer, 1909; Schölkopf &  
1209 Smola, 2002; Horn & Johnson, 2012).1210  
1211 **Symmetry.** Both  $(\mathbf{x} \cdot \mathbf{w})^2$  and  $\|\mathbf{x} - \mathbf{w}\|^2$  are symmetric in  $(\mathbf{x}, \mathbf{w})$ , hence  $k_{\mathbf{\Xi}}(\mathbf{x}, \mathbf{w}) = k_{\mathbf{\Xi}}(\mathbf{w}, \mathbf{x})$ .1212  
1213 **Factorization.** Write

1214 
$$k_{\mathbf{\Xi}}(\mathbf{x}, \mathbf{w}) = k_1(\mathbf{x}, \mathbf{w}) k_2(\mathbf{x}, \mathbf{w}), \quad k_1(\mathbf{x}, \mathbf{w}) = (\mathbf{x} \cdot \mathbf{w})^2, \quad k_2(\mathbf{x}, \mathbf{w}) = \frac{1}{\|\mathbf{x} - \mathbf{w}\|^2 + \varepsilon}.$$
  
1215  
1216

1217 *(a)  $k_1$  is PSD.* This is the homogeneous polynomial kernel of degree 2. With feature map  $\varphi(\mathbf{x}) =$   
1218  $\text{vec}(\mathbf{x}\mathbf{x}^\top)$  we have  $k_1(\mathbf{x}, \mathbf{w}) = \langle \varphi(\mathbf{x}), \varphi(\mathbf{w}) \rangle$ . Thus for any coefficients  $c_i$ ,

1219  
1220 
$$\sum_{i,j} c_i c_j k_1(\mathbf{x}_i, \mathbf{x}_j) = \left\| \sum_i c_i \varphi(\mathbf{x}_i) \right\|^2 \geq 0.$$
  
1221

(b)  $k_2$  is PSD (Gaussian Mixture). The inverse multiquadric kernel  $k_2$  admits a Laplace integral representation. Using the identity  $a^{-\beta} = \frac{1}{\Gamma(\beta)} \int_0^\infty t^{\beta-1} e^{-at} dt$  with  $\beta = 1$  and  $a = \|\mathbf{x} - \mathbf{w}\|^2 + \varepsilon$ , we have:

$$\frac{1}{\|\mathbf{x} - \mathbf{w}\|^2 + \varepsilon} = \int_0^\infty e^{-\varepsilon t} e^{-t\|\mathbf{x} - \mathbf{w}\|^2} dt.$$

For each fixed  $t > 0$ , the term  $e^{-t\|\mathbf{x} - \mathbf{w}\|^2}$  is a Gaussian kernel (RBF), which is known to be positive definite on  $\mathbb{R}^d$  (Buhmann, 2000). Since  $e^{-\varepsilon t} > 0$ ,  $k_2$  is a non-negative integral mixture of PD kernels, and is therefore PD (Fasshauer, 2011; dem, 1991).

(c) Product preserves PSD. The pointwise product of two PD kernels is PD (Schur product theorem). Since  $k_1$  and  $k_2$  are PD, their product  $k_{\mathbf{E}}$  is PD.

Thus  $k_{\mathbf{E}}$  is symmetric and PSD, hence a Mercer kernel on  $\mathbb{R}^d$ .  $\square$

### C.10 PROOF OF SELF-REGULATION FOR THE $\mathbf{E}$ -PRODUCT

**Proposition C.6** (The  $\mathbf{E}$ -Product is Naturally Self-Regulating). *For any fixed weight vector  $\mathbf{w}$ , the output of a  $\mathbf{E}$ -product neuron*

$$\mathbf{E}(\mathbf{w}, \mathbf{x}) = \frac{\langle \mathbf{w}, \mathbf{x} \rangle^2}{\|\mathbf{w} - \mathbf{x}\|^2 + \epsilon}$$

remains bounded and converges to a finite value as  $\|\mathbf{x}\| \rightarrow \infty$ .

*Proof.* Let  $\mathbf{x} = k\mathbf{u}$  where  $k = \|\mathbf{x}\|$  and  $\mathbf{u}$  is a unit vector. Then

$$\begin{aligned} \mathbf{E}(\mathbf{w}, k\mathbf{u}) &= \frac{\langle \mathbf{w}, k\mathbf{u} \rangle^2}{\|\mathbf{w} - k\mathbf{u}\|^2 + \epsilon} \\ &= \frac{k^2 \langle \mathbf{w}, \mathbf{u} \rangle^2}{\|\mathbf{w}\|^2 - 2k\langle \mathbf{w}, \mathbf{u} \rangle + k^2 + \epsilon}. \end{aligned}$$

Dividing numerator and denominator by  $k^2$  yields

$$\mathbf{E}(\mathbf{w}, k\mathbf{u}) = \frac{\langle \mathbf{w}, \mathbf{u} \rangle^2}{\frac{\|\mathbf{w}\|^2}{k^2} - \frac{2\langle \mathbf{w}, \mathbf{u} \rangle}{k} + 1 + \frac{\epsilon}{k^2}}.$$

Taking  $k \rightarrow \infty$ , all terms with  $k$  in the denominator vanish, and hence

$$\lim_{k \rightarrow \infty} \mathbf{E}(\mathbf{w}, k\mathbf{u}) = \langle \mathbf{w}, \mathbf{u} \rangle^2.$$

Since  $\langle \mathbf{w}, \mathbf{u} \rangle = \|\mathbf{w}\| \cos \theta$  for some angle  $\theta$ , the limit is

$$\|\mathbf{w}\|^2 \cos^2 \theta,$$

which lies in  $[0, \|\mathbf{w}\|^2]$ . Thus the  $\mathbf{E}$ -product output is bounded and convergent.

Note that this boundedness holds with respect to the inputs  $\mathbf{x}$ . To ensure the output remains bounded with respect to the weights and to support the self-regulation claim, we apply \*\*Weight Normalization\*\* (Salimans & Kingma, 2016). This decouples the magnitude of the weight vector from its direction, preventing quadratic growth of the output due to weight scaling.  $\square$

**Corollary C.7** (Dimensional Self-Normalization at Initialization). *Let  $\mathbf{x}, \mathbf{w} \in \mathbb{R}^d$  have i.i.d. components with zero mean and constant variance at initialization. Define*

$$A(\mathbf{x}, \mathbf{w}) = (\mathbf{x}^\top \mathbf{w})^2, \quad r(\mathbf{x}, \mathbf{w}) = \|\mathbf{x} - \mathbf{w}\|^2, \quad K(\mathbf{x}, \mathbf{w}) = \frac{A(\mathbf{x}, \mathbf{w})}{r(\mathbf{x}, \mathbf{w}) + \epsilon}.$$

1269 Then, as the dimension  $d \rightarrow \infty$ , both  $A(\mathbf{x}, \mathbf{w})$  and  $r(\mathbf{x}, \mathbf{w})$  scale as  $\mathcal{O}(d)$  in expectation, and the  
 1270 kernel value remains  $\mathcal{O}(1)$ :

$$1272 \quad \mathbb{E}[A(\mathbf{x}, \mathbf{w})] = \mathcal{O}(d), \quad \mathbb{E}[r(\mathbf{x}, \mathbf{w})] = \mathcal{O}(d), \quad \mathbb{E}[K(\mathbf{x}, \mathbf{w})] = \mathcal{O}(1).$$

1273 In particular, the  $\mathbf{E}$ -product is dimensionally self-normalizing at initialization.

1275 Proof. Write  $\mathbf{x} = (x_1, \dots, x_d)$  and  $\mathbf{w} = (w_1, \dots, w_d)$  with i.i.d. coordinates of zero mean and  
 1276 constant variance, and let  $\sigma_x^2 = \mathbb{E}[x_1^2]$ ,  $\sigma_w^2 = \mathbb{E}[w_1^2]$ . Then  
 1277

$$1278 \quad r(\mathbf{x}, \mathbf{w}) = \sum_{i=1}^d (x_i - w_i)^2 = \|\mathbf{x}\|^2 + \|\mathbf{w}\|^2 - 2\mathbf{x}^\top \mathbf{w}.$$

1281 By linearity of expectation and independence,

$$1282 \quad \mathbb{E}[\|\mathbf{x}\|^2] = d\sigma_x^2, \quad \mathbb{E}[\|\mathbf{w}\|^2] = d\sigma_w^2,$$

1284 while  $\mathbb{E}[\mathbf{x}^\top \mathbf{w}] = 0$  and  $\text{Var}(\mathbf{x}^\top \mathbf{w}) = \mathcal{O}(d)$ , so a typical realization of  $\mathbf{x}^\top \mathbf{w}$  has magnitude  $\mathcal{O}(\sqrt{d})$ .  
 1285 Hence

$$1286 \quad \mathbb{E}[r(\mathbf{x}, \mathbf{w})] = d(\sigma_x^2 + \sigma_w^2) + \mathcal{O}(\sqrt{d}) = \mathcal{O}(d).$$

1287 Similarly,  $\mathbf{x}^\top \mathbf{w} = \sum_{i=1}^d x_i w_i$  is a sum of  $d$  i.i.d. zero-mean variables with variance  $\mathcal{O}(1)$ , so  $\mathbf{x}^\top \mathbf{w}$   
 1288 has typical magnitude  $\mathcal{O}(\sqrt{d})$  and  
 1289

$$1290 \quad \mathbb{E}[A(\mathbf{x}, \mathbf{w})] = \mathbb{E}[(\mathbf{x}^\top \mathbf{w})^2] = \text{Var}(\mathbf{x}^\top \mathbf{w}) = \mathcal{O}(d).$$

1292 Combining these, we obtain the scaling

$$1293 \quad \mathbb{E}[K(\mathbf{x}, \mathbf{w})] = \mathbb{E}\left[\frac{A(\mathbf{x}, \mathbf{w})}{r(\mathbf{x}, \mathbf{w}) + \epsilon}\right] \approx \frac{\mathcal{O}(d)}{\mathcal{O}(d)} = \mathcal{O}(1),$$

1296 so the  $\mathbf{E}$ -product remains on a constant scale as  $d$  grows. This provides a heuristic scaling analysis  
 1297 showing that numerator and denominator are coupled in high dimensions, yielding a dimensionally  
 1298 self-normalizing kernel.  $\square$   
 1299

### 1300 C.11 ADDRESSING INTERNAL COVARIATE SHIFT

1302 **Corollary C.8** (Asymptotic Independence of Score Statistics). Let  $a = \mathbf{E}(\mathbf{w}, \mathbf{x})$  be the score of  
 1303 a neuron with weight vector  $\mathbf{w}$ . For a mini-batch  $\mathcal{B} = \{\mathbf{x}_1, \dots, \mathbf{x}_N\}$  where  $\mathbf{x}_i = k_i \mathbf{u}_i$ , define the  
 1304 empirical mean and variance:

$$1305 \quad \mu_{\mathcal{B}}(a) = \frac{1}{N} \sum_{i=1}^N a_i, \quad \sigma_{\mathcal{B}}^2(a) = \frac{1}{N} \sum_{i=1}^N (a_i - \mu_{\mathcal{B}}(a))^2.$$

1308 Then, as all  $k_i \rightarrow \infty$ ,

$$1310 \quad \lim_{k_1, \dots, k_N \rightarrow \infty} \mu_{\mathcal{B}}(a) = \|\mathbf{w}\|^2 \mathbb{E}_{\mathbf{u} \in \mathcal{U}} [\cos^2 \theta(\mathbf{w}, \mathbf{u})],$$

$$1312 \quad \lim_{k_1, \dots, k_N \rightarrow \infty} \sigma_{\mathcal{B}}^2(a) = \|\mathbf{w}\|^4 \text{Var}_{\mathbf{u} \in \mathcal{U}} [\cos^2 \theta(\mathbf{w}, \mathbf{u})],$$

1314 where  $\mathcal{U} = \{\mathbf{u}_1, \dots, \mathbf{u}_N\}$  is the set of normalized directions. Thus, asymptotically, the score statistics  
 1315 are independent of input magnitudes, mitigating internal covariate shift.

1316 *Proof.* From Proposition C.6, for  $\mathbf{x} = k\mathbf{u}$  with  $k \rightarrow \infty$ ,

$$\mathbb{E}(\mathbf{w}, \mathbf{x}) \rightarrow \|\mathbf{w}\|^2 \cos^2 \theta.$$

1319 To make this precise, fix  $\eta > 0$ . By Proposition C.6, for each  $i$  there exists  $K_i$  such that for all  
1320  $k_i \geq K_i$ ,

$$|a_i - \|\mathbf{w}\|^2 \cos^2 \theta_i| < \eta.$$

1323 Let  $K = \max_i K_i$ ; then for all  $k_i \geq K$  we have

$$1324 \quad \left| \mu_{\mathcal{B}}(a) - \|\mathbf{w}\|^2 \frac{1}{N} \sum_{i=1}^N \cos^2 \theta_i \right| \leq \eta.$$

1327 An analogous estimate, using  $|(a_i - \mu_{\mathcal{B}}(a))^2 - (\|\mathbf{w}\|^2 \cos^2 \theta_i - m)^2| \leq C\eta$  with  $m = \|\mathbf{w}\|^2 \frac{1}{N} \sum \cos^2 \theta_i$   
1328 and a constant  $C$  depending only on  $\|\mathbf{w}\|$  and the batch size, yields the variance limit. This proves  
1329 the stated limits for mean and variance. In particular, both statistics depend only on  $\|\mathbf{w}\|$  and the  
1330 angular distribution  $\mathcal{U}$ , not on magnitudes  $\{k_i\}$ , verifying mitigation of internal covariate shift.  $\square$

## 1332 C.12 PROOF OF STABLE LEARNING FOR THE $\mathbb{E}$ -PRODUCT

1334 **Proposition C.9** (The  $\mathbb{E}$ -Product Ensures Stable Learning). *The gradient of the  $\mathbb{E}$ -product with  
1335 respect to its input,  $\nabla_{\mathbf{x}} \mathbb{E}(\mathbf{w}, \mathbf{x})$ , approaches zero as the input vector  $\mathbf{x}$  moves infinitely far from the  
1336 weight vector  $\mathbf{w}$ .*

1338 *Proof.* We aim to prove that the learning signal, represented by the gradient of the  $\mathbb{E}$ -product with  
1339 respect to the input  $\mathbf{x}$ , diminishes for inputs that are distant from the learned weight vector  $\mathbf{w}$ .  
1340 This ensures that outliers do not cause large, destabilizing updates.

1341 The  $\mathbb{E}$ -product is defined as:

$$1343 \quad \mathbb{E}(\mathbf{w}, \mathbf{x}) = \frac{\langle \mathbf{w}, \mathbf{x} \rangle^2}{\|\mathbf{w} - \mathbf{x}\|^2 + \epsilon} = \frac{N(\mathbf{x})}{D(\mathbf{x})}$$

1346 where  $N(\mathbf{x}) = \langle \mathbf{w}, \mathbf{x} \rangle^2$  and  $D(\mathbf{x}) = \|\mathbf{w} - \mathbf{x}\|^2 + \epsilon$ .

1347 Using the quotient rule for vector calculus, the gradient  $\nabla_{\mathbf{x}} \mathbb{E}$  is:

$$1349 \quad \nabla_{\mathbf{x}} \mathbb{E} = \frac{(\nabla_{\mathbf{x}} N)D - N(\nabla_{\mathbf{x}} D)}{D^2}$$

1352 First, we compute the gradients of the numerator  $N(\mathbf{x})$  and the denominator  $D(\mathbf{x})$ :

### 1354 1. GRADIENT OF THE NUMERATOR

$$1356 \quad N(\mathbf{x}) = (\mathbf{w}^T \mathbf{x})^2$$

$$1357 \quad \nabla_{\mathbf{x}} N(\mathbf{x}) = 2(\mathbf{w}^T \mathbf{x}) \cdot \nabla_{\mathbf{x}} (\mathbf{w}^T \mathbf{x}) = 2\langle \mathbf{w}, \mathbf{x} \rangle \mathbf{w}$$

### 1359 2. GRADIENT OF THE DENOMINATOR

$$1360 \quad D(\mathbf{x}) = \|\mathbf{w} - \mathbf{x}\|^2 + \epsilon = (\mathbf{w} - \mathbf{x})^T (\mathbf{w} - \mathbf{x}) + \epsilon$$

$$1362 \quad \nabla_{\mathbf{x}} D(\mathbf{x}) = 2(\mathbf{w} - \mathbf{x}) \cdot (-1) = -2(\mathbf{w} - \mathbf{x}) = 2(\mathbf{x} - \mathbf{w})$$

1363 Substituting these into the quotient rule expression:

$$\nabla_{\mathbf{x}} \mathbb{E} = \frac{(2\langle \mathbf{w}, \mathbf{x} \rangle \mathbf{w})(\|\mathbf{w} - \mathbf{x}\|^2 + \epsilon) - (\langle \mathbf{w}, \mathbf{x} \rangle^2)(2(\mathbf{x} - \mathbf{w}))}{(\|\mathbf{w} - \mathbf{x}\|^2 + \epsilon)^2}$$

1368 To analyze the behavior for distant inputs, we examine the limit as  $\|\mathbf{x}\| \rightarrow \infty$ . Let  $\mathbf{x} = k\mathbf{u}$ , where  
1369  $k = \|\mathbf{x}\|$  and  $\mathbf{u}$  is a unit vector.

1370 As  $k \rightarrow \infty$ :

- 1372 •  $\langle \mathbf{w}, \mathbf{x} \rangle = k\langle \mathbf{w}, \mathbf{u} \rangle \sim \mathcal{O}(k)$
- 1373 •  $\|\mathbf{w} - \mathbf{x}\|^2 = \|\mathbf{w}\|^2 - 2k\langle \mathbf{w}, \mathbf{u} \rangle + k^2 \sim \mathcal{O}(k^2)$

1376 Let's analyze the order of magnitude for the terms in the gradient's numerator:

- 1377 • First term:  $(2\langle \mathbf{w}, \mathbf{x} \rangle \mathbf{w})(\|\mathbf{w} - \mathbf{x}\|^2 + \epsilon) \sim \mathcal{O}(k) \cdot \mathcal{O}(k^2) = \mathcal{O}(k^3)$
- 1378 • Second term:  $(\langle \mathbf{w}, \mathbf{x} \rangle^2)(2(\mathbf{x} - \mathbf{w})) \sim \mathcal{O}(k^2) \cdot \mathcal{O}(k) = \mathcal{O}(k^3)$

1381 The numerator as a whole is of order  $\mathcal{O}(k^3)$ .

1382 The denominator is  $(\|\mathbf{w} - \mathbf{x}\|^2 + \epsilon)^2 \sim (\mathcal{O}(k^2))^2 = \mathcal{O}(k^4)$ .

1384 Therefore, the magnitude of the gradient behaves as:

$$1386 \quad \|\nabla_{\mathbf{x}} \mathbb{E}\| \sim \frac{\mathcal{O}(k^3)}{\mathcal{O}(k^4)} = \mathcal{O}\left(\frac{1}{k}\right)$$

1388 As  $k = \|\mathbf{x}\| \rightarrow \infty$ , the magnitude of the gradient approaches zero:

$$1390 \quad \lim_{\|\mathbf{x}\| \rightarrow \infty} \|\nabla_{\mathbf{x}} \mathbb{E}(\mathbf{w}, \mathbf{x})\| = 0$$

1393 This proves that for inputs  $\mathbf{x}$  that are very far from the weight vector  $\mathbf{w}$ , the gradient becomes  
1394 vanishingly small. While this provides robustness against outliers, it also implies a risk of van-  
1395 ishing gradients at initialization if weights are not properly scaled. To ensure stability and avoid  
1396 barren plateaus, we mandate **\*\*Data Normalization\*\*** and **\*\*Weight Normalization\*\*** (Salimans  
1397 & Kingma, 2016). Specifically, initializing with normalized random weights ensures the network  
1398 starts in an active regime. Note that unlike standard RBFs, we do not require negative weights or  
1399 complex initialization schemes because the superposition of the weight vectors (via the quadratic  
1400 numerator) naturally covers the space.  $\square$

### 1401 C.13 REGULARITY PROPERTIES OF THE $\mathbb{E}$ -PRODUCT KERNEL

1403 **Lemma C.10** (Analyticity). *For any fixed  $\mathbf{w} \in \mathbb{R}^d$  and  $\epsilon > 0$ , the map*

$$1404 \quad \mathbf{x} \mapsto K(\mathbf{w}, \mathbf{x})$$

1406 *is real-analytic on  $\mathbb{R}^d$ , and in particular infinitely differentiable ( $C^\infty$ ).*

1408 *Proof.* Both  $N(\mathbf{x}) = \langle \mathbf{w}, \mathbf{x} \rangle^2$  and  $D(\mathbf{x}) = \|\mathbf{w} - \mathbf{x}\|^2 + \epsilon$  are polynomials; since  $D \geq \epsilon > 0$ ,  $1/D$  is  
1409 analytic and so is  $K = N/D$ .  $\square$

1410  
1411   **Proposition C.11** (Lipschitz Continuity). *For any fixed  $\mathbf{w} \in \mathbb{R}^d$  and  $\epsilon > 0$ , the kernel  $K(\mathbf{w}, \mathbf{x})$  is*  
1412   *globally Lipschitz continuous in  $\mathbf{x}$ .*

1413   *Proof.* It suffices to show that  $\sup_{\mathbf{x} \in \mathbb{R}^d} \|\nabla_{\mathbf{x}} K(\mathbf{w}, \mathbf{x})\| < \infty$ .

1415   Writing  $K = N/D$  with  $N(\mathbf{x}) = \langle \mathbf{w}, \mathbf{x} \rangle^2$  and  $D(\mathbf{x}) = \|\mathbf{w} - \mathbf{x}\|^2 + \epsilon$ , we compute

$$\nabla_{\mathbf{x}} K(\mathbf{w}, \mathbf{x}) = \frac{2\langle \mathbf{w}, \mathbf{x} \rangle \mathbf{w} D(\mathbf{x}) - 2\langle \mathbf{w}, \mathbf{x} \rangle^2 (\mathbf{x} - \mathbf{w})}{D(\mathbf{x})^2}.$$

1419   Let  $a = \|\mathbf{w}\|$ ,  $r = \|\mathbf{x}\|$ . By Cauchy–Schwarz,

$$|\langle \mathbf{w}, \mathbf{x} \rangle| \leq ar, \quad \|\mathbf{x} - \mathbf{w}\| \leq r + a.$$

1420   Moreover

$$D(\mathbf{x}) = \|\mathbf{w} - \mathbf{x}\|^2 + \epsilon \geq (r - a)^2 + \epsilon.$$

1421   Substituting these bounds gives

$$\|\nabla_{\mathbf{x}} K(\mathbf{w}, \mathbf{x})\| \leq \frac{2a^2r}{(r - a)^2 + \epsilon} + \frac{2a^2r^2(r + a)}{\left((r - a)^2 + \epsilon\right)^2} =: \Phi(r).$$

1422   As  $r \rightarrow \infty$ ,  $\Phi(r) = O(1/r) \rightarrow 0$ . Since  $\Phi$  is continuous on  $[0, \infty)$  and bounded at infinity, it attains  
1423   a finite global maximum:

$$L(a, \epsilon) := \sup_{r \geq 0} \Phi(r) < \infty.$$

1432   By the mean value theorem in  $\mathbb{R}^d$ ,

$$|K(\mathbf{w}, \mathbf{x}) - K(\mathbf{w}, \mathbf{y})| \leq L(a, \epsilon) \|\mathbf{x} - \mathbf{y}\|,$$

1435   so  $K$  is globally Lipschitz in  $\mathbf{x}$  with constant  $L(a, \epsilon)$ , depending on  $\|\mathbf{w}\|$  and  $\epsilon$ .  $\square$

## 1437   C.14   UNIVERSAL APPROXIMATION THEOREM FOR $\mathbb{E}$ -PRODUCT NETWORKS

### 1438   SETUP AND NOTATION

1440   Let  $\mathcal{X} \subset \mathbb{R}^d$  be compact and fix  $\varepsilon > 0$ . Define the unit

$$g(x; w, b) := \frac{(x \cdot w + b)^2}{\|x - w\|^2 + \varepsilon}, \quad w \in \mathbb{R}^d, b \in \mathbb{R},$$

1444   and let

$$\mathcal{F} = \text{span}\{g(\cdot; w, b) \mid w \in \mathbb{R}^d, b \in \mathbb{R}\}.$$

1446   Denote the inverse–multiquadric (IMQ) kernel centered at  $w$  by

$$K(\cdot, w) := (\|\cdot - w\|^2 + \varepsilon)^{-1}.$$

1448   **Lemma C.12** (Derivatives in the bias recover IMQ). *For every fixed  $w \in \mathbb{R}^d$  and  $b \in \mathbb{R}$  there*  
1449   *holds, pointwise in  $x$ ,*

$$\frac{\partial^2}{\partial b^2} g(x; w, b) = \frac{2}{\|x - w\|^2 + \varepsilon}.$$

1451   *Moreover, for compact  $\mathcal{X}$  the second derivative in  $b$  is the uniform limit on  $\mathcal{X}$  of finite difference*  
1452   *quotients built from translates of  $g$ , and hence*

$$K(\cdot, w) = \frac{1}{2} \partial_b^2 g(\cdot; w, b) \in \overline{\mathcal{F}}$$

1455   *(the uniform closure on  $\mathcal{X}$ ).*

1457 *Proof.* Fix  $x, w$  and write

$$1458 \quad g(x; w, b) = \frac{(x \cdot w)^2 + 2b(x \cdot w) + b^2}{\|x - w\|^2 + \varepsilon}.$$

1460 Differentiating in  $b$  yields

$$1462 \quad \partial_b g(x; w, b) = \frac{2(x \cdot w) + 2b}{\|x - w\|^2 + \varepsilon}, \quad \partial_b^2 g(x; w, b) = \frac{2}{\|x - w\|^2 + \varepsilon},$$

1464 establishing the first identity.

1466 Next, for each fixed  $x \in \mathcal{X}$  the function  $b \mapsto g(x; w, b)$  is a polynomial in  $b$  whose denominator  
1467 satisfies  $\|x - w\|^2 + \varepsilon \geq \varepsilon > 0$ . Hence  $g$  is  $C^\infty$  in  $b$  and all derivatives  $\partial_b^k g(\cdot; w, b)$  are continuous  
1468 and uniformly bounded on  $\mathcal{X}$ . Thus finite difference quotients for  $\partial_b$  and  $\partial_b^2$  converge uniformly on  
1469  $\mathcal{X}$  to the corresponding derivatives.

1470 Each finite-difference quotient is a linear combination of finitely many translates  $g(\cdot; w, b + h)$  and  
1471  $g(\cdot; w, b)$ , hence belongs to  $\mathcal{F}$ . Passing to the uniform limit shows that  $\partial_b^2 g(\cdot; w, b) \in \overline{\mathcal{F}}$ . Substituting  
1472 the formula above yields

$$1473 \quad K(\cdot, w) = \frac{1}{2} \partial_b^2 g(\cdot; w, b) \in \overline{\mathcal{F}}.$$

1474  $\square$

1475 **Lemma C.13** (Fourier transform of IMQ and positivity). *Let  $K(x) = (\|x\|^2 + \varepsilon)^{-1}$  on  $\mathbb{R}^d$  with  
1476  $\varepsilon > 0$ . Its Fourier transform (distributional for  $d \geq 4$ ; classical otherwise) is radial and equals*

$$1478 \quad \widehat{K}(\xi) = C_{d,\varepsilon} \|\xi\|^{1-d/2} K_{\frac{d}{2}-1}(\sqrt{\varepsilon} \|\xi\|), \quad (28)$$

1480 where  $K_\nu$  is the modified Bessel function of the second kind and  $C_{d,\varepsilon} > 0$ . In particular,  $\widehat{K}(\xi) > 0$   
1481 for all  $\xi \neq 0$ , with the removable value at  $\xi = 0$  taken by continuity.

1482 *Proof.* Because  $K$  is radial, its Fourier transform is radial and given by the Hankel transform  
1483 formula

$$1485 \quad \widehat{K}(\xi) = (2\pi)^{d/2} \|\xi\|^{-(d/2-1)} \int_0^\infty r^{d/2} K(r) J_{d/2-1}(r \|\xi\|) dr.$$

1487 Use the Laplace representation

$$1489 \quad K(r) = \frac{1}{r^2 + \varepsilon} = \int_0^\infty e^{-t(r^2 + \varepsilon)} dt,$$

1491 valid by Tonelli since the integrand is positive. Substitute and switch integrals:

$$1493 \quad \widehat{K}(\xi) = (2\pi)^{d/2} \|\xi\|^{-(d/2-1)} \int_0^\infty e^{-t\varepsilon} \left( \int_0^\infty r^{d/2} e^{-tr^2} J_{d/2-1}(r \|\xi\|) dr \right) dt.$$

1495 The inner integral is standard:

$$1497 \quad \int_0^\infty r^{\nu+1} e^{-tr^2} J_\nu(rs) dr = \frac{s^\nu}{(2t)^{\nu+1}} e^{-s^2/(4t)}, \quad \nu > -1, t > 0.$$

1499 Applying this with  $\nu = d/2 - 1$  and simplifying gives an integral representation whose evaluation  
1500 is the modified Bessel function  $K_\nu$ . The resulting constants combine to give (28) with  $C_{d,\varepsilon} > 0$ .

1502 Since  $K_\nu(z) > 0$  for all  $z > 0$ , we obtain  $\widehat{K}(\xi) > 0$  for all  $\xi \neq 0$ . Continuity gives positivity at  $\xi = 0$   
1503 as well.  $\square$

1504  
 1505   **Theorem C.14** (Universal approximation for  $\mathbb{E}$ -product networks). *The class  $\mathcal{F}$  is dense in  $C(\mathcal{X})$*   
 1506   *under the uniform norm.*

1507   *Proof.* Lemma C.12 implies

$$1508 \quad K(\cdot, w) \in \overline{\mathcal{F}} \quad \text{for every } w \in \mathbb{R}^d.$$

1509   Hence

$$1511 \quad \text{span}\{K(\cdot, w) : w \in \mathbb{R}^d\} \subseteq \overline{\mathcal{F}}.$$

1512   Let  $\mu$  be a finite signed Borel measure supported on  $\mathcal{X}$  such that

$$1514 \quad \int_{\mathcal{X}} K(x, w) d\mu(x) = 0 \quad \forall w \in \mathbb{R}^d.$$

1516   Extend  $\mu$  by 0 outside  $\mathcal{X}$  to obtain a compactly supported measure on  $\mathbb{R}^d$ . The integral condition  
 1517   is equivalent to

$$1518 \quad (K * \mu)(w) = 0, \quad \forall w \in \mathbb{R}^d.$$

1520   Taking Fourier transforms gives

$$1521 \quad \widehat{K}(\xi) \widehat{\mu}(\xi) = 0.$$

1522   By Lemma C.13,  $\widehat{K}(\xi) > 0$  for all  $\xi$ , hence  $\widehat{\mu}(\xi) \equiv 0$ . Uniqueness of Fourier transform for compactly  
 1523   supported finite measures implies  $\mu \equiv 0$ .

1524   By the Hahn–Banach / Riesz duality criterion, the span of  $\{K(\cdot, w) : w \in \mathbb{R}^d\}$  is dense in  $C(\mathcal{X})$ .  
 1525   Since this span is contained in  $\overline{\mathcal{F}}$ , we conclude  $\overline{\mathcal{F}} = C(\mathcal{X})$ .  $\square$

1527   **Corollary C.15** (Practical corollary: single-hidden-layer networks are universal). *Let networks of  
 1528   the form*

$$1529 \quad x \mapsto \sum_{i=1}^n \alpha_i g(x; w_i, b_i) + c$$

1532   be considered. Then for every continuous  $f \in C(\mathcal{X})$  and every  $\delta > 0$  there exist  $n$  and parameters  
 1533    $\{\alpha_i, w_i, b_i, c\}$  such that

$$1534 \quad \sup_{x \in \mathcal{X}} \left| f(x) - \sum_{i=1}^n \alpha_i g(x; w_i, b_i) - c \right| < \delta.$$

1537   *Proof.* Immediate from Theorem C.14.  $\square$

## 1539   C.15 INFORMATION-GEOMETRIC FOUNDATIONS OF THE $\mathbb{E}$ -PRODUCT

### 1541   C.15.1 DEFINITION AND GEOMETRIC INTERPRETATION

1542   We consider probability distributions in the simplex  $\Delta^{n-1} = \{\mathbf{p} \in \mathbb{R}_{\geq 0}^n : \sum_{i=1}^n p_i = 1\}$ . While  
 1543   information geometry traditionally employs the Fisher metric, we establish a novel connection to  
 1544   Euclidean geometry through the  $\mathbb{E}$ -product.

1546   **Definition C.1** ( $\mathbb{E}$ -Product: Geometric Similarity Measure). *For distinct distributions  $\mathbf{p}, \mathbf{q} \in$   
 1547    $\Delta^{n-1}$ , the  $\mathbb{E}$ -product is defined as:*

$$1548 \quad \mathbb{E}(\mathbf{p}, \mathbf{q}) := \frac{(\mathbf{p} \cdot \mathbf{q})^2}{\|\mathbf{p} - \mathbf{q}\|_2^2}$$

1550   *where:*

- 1551     •  $\mathbf{p} \cdot \mathbf{q} = \sum_{i=1}^n p_i q_i$  measures distributional alignment  
 1552     •  $\|\mathbf{p} - \mathbf{q}\|_2^2 = \sum_{i=1}^n (p_i - q_i)^2$  quantifies Euclidean dissimilarity

1553     This ratio captures the tension between distributional agreement and geometric separation.

1554     **Remark C.16** (Singularity and Invariance Properties). When  $\mathbf{p} = \mathbf{q}$ , we define  $\mathbb{E}(\mathbf{p}, \mathbf{q}) := \infty$  via  
 1555     the limit:

$$\lim_{\mathbf{q} \rightarrow \mathbf{p}} \mathbb{E}(\mathbf{p}, \mathbf{q}) = \infty$$

1556     reflecting maximal self-similarity. The  $\mathbb{E}$ -product exhibits two key properties:

- 1557       1. **Symmetry:**  $\mathbb{E}(\mathbf{p}, \mathbf{q}) = \mathbb{E}(\mathbf{q}, \mathbf{p})$   
 1558       2. **Scale Invariance:** Invariant under index permutation

### 1559     C.15.2 EXTREMAL SIMILARITY THEOREMS

1560     We collect here the detailed proofs of the extremal similarity results stated in the main text.

#### 1561     Proof of Minimal Similarity and Statistical Orthogonality.

1562     *Proof.* ( $\Rightarrow$ ) Assume  $\mathbb{E}(\mathbf{p}, \mathbf{q}) = 0$ . Since  $\mathbf{p} \neq \mathbf{q}$ ,  $\|\mathbf{p} - \mathbf{q}\|_2^2 > 0$ . Thus  $(\mathbf{p} \cdot \mathbf{q})^2 = 0 \Rightarrow \sum p_i q_i = 0$ . By  
 1563     non-negativity of probabilities,  $p_i q_i = 0 \forall i$ , hence  $\text{supp}(\mathbf{p}) \cap \text{supp}(\mathbf{q}) = \emptyset$ .

1564     ( $\Leftarrow$ ) Disjoint supports imply  $\forall i : (p_i > 0 \Rightarrow q_i = 0)$  and vice versa. Thus  $\mathbf{p} \cdot \mathbf{q} = 0$ , so  $\mathbb{E}(\mathbf{p}, \mathbf{q}) = 0$ .

1565     The KL divergence  $\text{KL}(\mathbf{p} \parallel \mathbf{q})$  contains terms  $\log(p_i/q_i)$  where  $p_i > 0$  and  $q_i = 0$ , causing divergence.  
 1566     Similar reasoning applies to  $\text{KL}(\mathbf{q} \parallel \mathbf{p})$  and cross-entropy  $H(\mathbf{p}, \mathbf{q})$  (Cover & Thomas, 2006).  $\square$

#### 1567     Proof of Maximal Similarity and Distributional Identity.

1568     *Proof.* ( $\Rightarrow$ ) Suppose  $\mathbb{E}(\mathbf{p}, \mathbf{q}) \rightarrow \infty$ . By Cauchy-Schwarz (Horn & Johnson, 2012),  $\mathbf{p} \cdot \mathbf{q} \leq \|\mathbf{p}\|_2 \|\mathbf{q}\|_2 \leq 1$ . Since the numerator is bounded,  $\|\mathbf{p} - \mathbf{q}\|_2^2 \rightarrow 0$ , implying  $\mathbf{p} = \mathbf{q}$ .

1569     ( $\Leftarrow$ ) For  $\mathbf{p} = \mathbf{q}$ , consider  $\mathbf{q}^{(k)} \rightarrow \mathbf{p}$ . Then:

$$\mathbf{p} \cdot \mathbf{q}^{(k)} \rightarrow \|\mathbf{p}\|_2^2 \geq \frac{1}{n} > 0 \quad (\text{since } \|\mathbf{p}\|_2^2 \geq \frac{1}{n} \text{ by Cauchy-Schwarz})$$

1570     while  $\|\mathbf{p} - \mathbf{q}^{(k)}\|_2^2 \rightarrow 0$ , so  $\mathbb{E}(\mathbf{p}, \mathbf{q}^{(k)}) \rightarrow \infty$ .

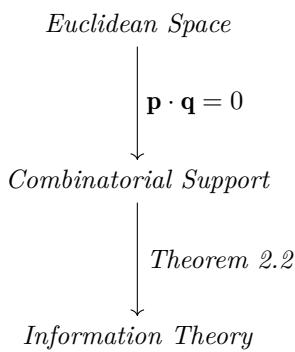
1571     When  $\mathbf{p} = \mathbf{q}$ ,  $\log(p_i/q_i) = 0$  for all  $i$ , so  $\text{KL}(\mathbf{p} \parallel \mathbf{q}) = 0$ . Cross-entropy reduces to entropy when  
 1572     distributions are identical.  $\square$

1573     **Remark C.17** (Duality of Orthogonality Concepts). The  $\mathbb{E}$ -product unifies three distinct notions  
 1574     of orthogonality:

$$\begin{aligned} \text{Euclidean: } \mathbf{p} \perp \mathbf{q} &\iff \mathbf{p} \cdot \mathbf{q} = 0 \\ \text{Combinatorial: } \text{supp}(\mathbf{p}) \cap \text{supp}(\mathbf{q}) &= \emptyset \\ \text{Information-Theoretic: } \text{KL}(\mathbf{p} \parallel \mathbf{q}) &= \infty \end{aligned}$$

1575     Theorem 2.2 establishes their equivalence through  $\mathbb{E}(\mathbf{p}, \mathbf{q}) = 0$ . This contrasts with Fisher-based  
 1576     orthogonality, which depends on manifold curvature.

1577     **Remark C.18** (Geometric-Information Duality). The  $\mathbb{E}$ -product creates a bridge between geometric  
 1578     and probabilistic perspectives:



## 1611 D DETAILED XOR ANALYSIS AND GRADIENT PROPERTIES

1613 This section provides a comprehensive analysis of the  $\mathbb{E}$ -product’s ability to solve the XOR problem  
1614 and its gradient properties that enable stable learning.

### 1616 D.1 MATHEMATICAL FOUNDATION FOR XOR SOLUTION

1618 The  $\mathbb{E}$ -product’s non-linearity is not merely a mathematical curiosity; it has practical implications  
1619 for neural computation. By integrating alignment and proximity into a single operator, the  $\mathbb{E}$ -  
1620 product allows for more nuanced feature learning. It can adaptively respond to inputs based on  
1621 their geometric relationships with learned weight vectors, enabling the network to capture complex  
1622 patterns without the need for separate activation functions.

1623 The XOR problem is not linearly separable and thus cannot be solved by a single traditional neuron  
1624 (linear perceptron). However, a single  $\mathbb{E}$ -product unit can solve this problem due to its inherent  
1625 non-linearity. We have formally proven that the  $\mathbb{E}$ -product is a valid Mercer kernel (Theorem 2.1;  
1626 see Appendix C.9)(Mercer, 1909; Hofmann et al., 2008; Micchelli et al., 2006; Schölkopf et al., 1997;  
1627 Cortes, 1995; Jacot et al., 2018).

### 1628 D.2 GRADIENT ANALYSIS AND LEARNING STABILITY

1630 To understand the  $\mathbb{E}$ -product’s behavior during learning, we analyze its gradient properties. A key  
1631 property for stable training is that the gradient with respect to the input,  $\nabla_x \mathcal{K}_{\mathbb{E}}$ , diminishes as the  
1632 input  $x$  moves far from the weight vector  $w$ . This ensures that distant outliers do not cause large,  
1633 destabilizing updates. We have formally proven this as a stability Proposition (Proposition C.9;  
1634 Appendix C.12), demonstrating that  $\lim_{\|x\| \rightarrow \infty} \|\nabla_x \mathcal{K}_{\mathbb{E}}(w, x)\| = 0$ .

1635 The presence of  $\epsilon$  in the denominator ensures that the derivative remains well-defined, avoiding  
1636 division by zero and contributing to numerical stability. This contrasts with activation functions  
1637 like ReLU, which have a derivative of zero for negative inputs, potentially leading to ”dead neurons.”  
1638 The smooth and generally non-zero gradient of the  $\mathbb{E}$ -product is hypothesized to contribute to more  
1639 stable and efficient learning dynamics, reducing the reliance on auxiliary mechanisms like complex  
1640 normalization schemes.

1641 The non-linearity is thus not an add-on but an intrinsic property derived from the direct mathemati-  
1642 cal interaction of vector projection (alignment, via the  $(w^T x)^2$  term) and vector distance (proximity,  
1643 via the  $\|w - x\|^2 + \epsilon$  term). This provides a mathematically grounded basis for feature learning, as  
1644 the unit becomes selectively responsive to inputs that exhibit specific geometric relationships, both

1645 in terms of angular alignment and spatial proximity, to its learned weight vector  $\mathbf{w}$ . Consequently,  
 1646  $\mathbf{w}$  can be interpreted as a learned feature template or prototype that the unit is tuned to detect,  
 1647 with the  $\mathbb{E}$ -product quantifying the degree of match in a nuanced, non-linear fashion.  
 1648

### 1649 D.3 OPTIMIZATION LANDSCAPE ANALYSIS 1650

1651 The gradient of the  $\mathbb{E}$ -product, being responsive across the input space, actively pushes the neuron’s  
 1652 weights away from configurations that would lead to a zero output (neuron death, e.g., at an input  
 1653 of  $[0, 0]$  for this problem if weights were also near zero). This contrasts with a simple dot product  
 1654 neuron where the gradient might vanish or lead to a global minimum at zero output for certain  
 1655 problems.

1656 For instance, when considering gradient-based optimization, the loss landscape “seen” by the  $\mathbb{E}$ -  
 1657 product neuron in the XOR context would exhibit a peak or high loss at  $[0, 0]$  (if that were the  
 1658 target for non-zero outputs), encouraging weights to move towards a state that correctly classifies.  
 1659 Conversely, a simple dot product neuron might present a loss landscape where a gradient-based  
 1660 optimizer could find a stable (but incorrect) minimum at zero output. This ability to avoid such  
 1661 dead zones and actively shape the decision boundary makes it helpful to solve problems like XOR  
 1662 with a single unit, leveraging its inherent non-linearity as a mathematical kernel.

### 1663 D.4 GEOMETRIC INTERPRETATION OF DECISION BOUNDARIES 1664

1665 Conceptually, the decision boundary or vector field generated by a simple dot product neuron is  
 1666 linear, forming a hyperplane that attempts to separate data points. In contrast, the  $\mathbb{E}$ -product  
 1667 generates a more complex, non-linear vector field. This field can be visualized as creating a series  
 1668 of potential wells or peaks centered around the weight vector  $\mathbf{w}$ , with the strength of influence  
 1669 decaying with distance. This structure allows for more nuanced and localized responses. The  
 1670  $\mathbb{E}$ -product solves XOR through its Localized Quadratic Expressivity. The quadratic numerator  
 1671  $(\mathbf{w}^\top \mathbf{x})^2$  provides the necessary non-linearity, exhibiting a Superposition property that manifests as  
 1672 Sign Invariance: the response is identical for  $\mathbf{x}$  and  $-\mathbf{x}$ . This allows the network to classify antipodal  
 1673 points similarly. Geometrically, the solution corresponds to an orthogonal valley—the region where  
 1674  $\mathbf{w}^\top \mathbf{x} \approx 0$ —where the loss is minimal, effectively separating the XOR classes.

1675  
 1676 Figure 6: Comparison of loss landscapes: dot-product neuron (left) exhibits spurious minima leading  
 1677 to vanishing gradients, while  $\mathbb{E}$ -product neuron (right) yields non-degenerate gradients enabling  
 1678 optimization to reach correct separators.  
 1679

## 1680 E DETAILED VORTEX DYNAMICS ANALYSIS 1681

1682 This section provides a comprehensive mathematical analysis of the vortex-like learning dynamics  
 1683 exhibited by  $\mathbb{E}$ -product neurons and their space-partitioning behavior.  
 1684

### 1685 E.1 FUNDAMENTAL LEARNING DYNAMICS 1686

1687 We begin by analyzing the fundamental learning dynamics that emerge in both conventional and our  
 1688 proposed architectures. In artificial intelligence, competitive learning manifests in various forms,  
 1689 whether through linear classification using dot products or clustering using Euclidean distances.  
 1690 Both approaches involve partitioning the feature space between neurons, which can be conceptu-  
 1691

1692 alized as prototype learning where each neuron claims a territorial "field" in the representation  
 1693 space.  
 1694

## 1695 E.2 CONVENTIONAL VS. $\mathbb{E}$ -PRODUCT DECISION BOUNDARIES 1696

1697 In a conventional linear model, the logit for each class  $i$  is computed as:

$$1698 \quad z_i = \mathbf{w}_i^T \mathbf{x} \quad (29)$$

1700 where  $\mathbf{w}_i$  is the weight vector (prototype) for class  $i$ , and  $\mathbf{x}$  is the input vector. The softmax function  
 1701 then normalizes these logits into probabilities:

$$1702 \quad p_i = \frac{\exp(z_i)}{\sum_{j=1}^C \exp(z_j)} \quad (30)$$

1705 The decision boundary between any two classes  $i$  and  $j$  forms a linear hyperplane defined by:  
 1706

$$1707 \quad (\mathbf{w}_i - \mathbf{w}_j)^T \mathbf{x} = 0 \quad (31)$$

1708 During training via gradient descent, each prototype  $\mathbf{w}_i$  is updated to maximize its alignment with  
 1709 the data distributions of its assigned class. This optimization process often leads to an unbounded  
 1710 increase in prototype magnitudes, as  $\|\mathbf{w}_i\| \rightarrow \infty$  directly amplifies the logit  $z_i$ , thereby increasing  
 1711 the model's confidence. However, the decision boundaries themselves remain linear hyperplanes,  
 1712 creating rigid geometric separations in the feature space.

1713 In contrast, the non-linear  $\mathbb{E}$ -product allows neurons to learn more representative prototypes for  
 1714 each class, leading to the formation of more nuanced decision boundaries. For the  $\mathbb{E}$ -product, the  
 1715 response of neuron  $i$  to input  $\mathbf{x}$  is given by:

$$1717 \quad z_i = \mathbb{E}(\mathbf{w}_i, \mathbf{x}) = \frac{\langle \mathbf{w}_i, \mathbf{x} \rangle^2}{\|\mathbf{w}_i - \mathbf{x}\|^2 + \epsilon} \quad (32)$$

1719 This formulation embodies the signal-to-noise ratio interpretation established in our theoretical  
 1720 framework (Appendix C.15), where the squared dot product  $\langle \mathbf{w}_i, \mathbf{x} \rangle^2$  represents the "signal" of  
 1721 distributional alignment, and  $\|\mathbf{w}_i - \mathbf{x}\|^2$  quantifies the "noise" of dissimilarity.

## 1723 E.3 SOFTMAX DYNAMICS AND COMPETITIVE LEARNING 1724

1725 Similarly to conventional neurons, the  $\mathbb{E}$ -product outputs are normalized using the softmax function:

$$1727 \quad p_i = \frac{\exp(z_i)}{\sum_{j=1}^C \exp(z_j)} = \frac{\exp\left(\frac{\langle \mathbf{w}_i, \mathbf{x} \rangle^2}{\|\mathbf{w}_i - \mathbf{x}\|^2 + \epsilon}\right)}{\sum_{j=1}^C \exp\left(\frac{\langle \mathbf{w}_j, \mathbf{x} \rangle^2}{\|\mathbf{w}_j - \mathbf{x}\|^2 + \epsilon}\right)} \quad (33)$$

1731 This softmax normalization serves a crucial role in the competitive dynamics of  $\mathbb{E}$ -product neurons.  
 1732 The softmax function acts as a transformation that maps from the real-valued  $\mathbb{E}$ -product responses  
 1733  $z_i \in \mathbb{R}$  to a delta distribution  $\delta_i$  in probability space. This softmax distribution over  $\mathbb{E}$ -product  
 1734 scores can be interpreted as the *posterior responsibility* of each prototype (neuron) for the input,  
 1735 drawing a direct connection to Gaussian Mixture Models (GMMs) and expectation-maximization  
 1736 frameworks.

1737 The softmax can also be viewed as computing a categorical distribution proportional to exponentiated  
 1738 log-likelihoods, which in this case derive from a geometric  $\mathbb{E}$ -product similarity rather than

1739 traditional probabilistic assumptions. This bridges the gap between probabilistic views (such as  
 1740 EM algorithms and classification) and our geometric formulation, providing a principled foundation  
 1741 for the competitive dynamics.

1742 As training progresses and the differences between  $\mathbb{E}$ -product responses become more pronounced,  
 1743 the softmax transformation approaches a delta distribution, where the winning neuron (with the  
 1744 highest  $\mathbb{E}$ -product response) approaches probability 1 while all others approach 0. This yields  
 1745 competitive learning dynamics in which neurons specialize on distinct regions of the input space.  
 1746

#### 1747 E.4 MATHEMATICAL CHARACTERIZATION OF DECISION BOUNDARIES

1749 The decision boundary between two neurons with prototypes  $\mathbf{w}_i$  and  $\mathbf{w}_j$  is defined by the condition  
 1750 where their responses are equal:

$$1751 \quad \mathbb{E}(\mathbf{w}_i, \mathbf{x}) = \mathbb{E}(\mathbf{w}_j, \mathbf{x}) \quad (34)$$

1753 Expanding this condition:

$$1754 \quad \frac{\langle \mathbf{w}_i, \mathbf{x} \rangle^2}{\|\mathbf{w}_i - \mathbf{x}\|^2 + \epsilon} = \frac{\langle \mathbf{w}_j, \mathbf{x} \rangle^2}{\|\mathbf{w}_j - \mathbf{x}\|^2 + \epsilon} \quad (35)$$

1757 Cross-multiplying and rearranging:

$$1759 \quad \langle \mathbf{w}_i, \mathbf{x} \rangle^2 (\|\mathbf{w}_j - \mathbf{x}\|^2 + \epsilon) = \langle \mathbf{w}_j, \mathbf{x} \rangle^2 (\|\mathbf{w}_i - \mathbf{x}\|^2 + \epsilon) \quad (36)$$

1761 This equation defines a non-linear decision boundary that depends on both alignment (via squared  
 1762 dot products) and proximity (via squared distances) between the input and each prototype. Unlike  
 1763 the linear hyperplane formed by conventional dot-product neurons, the  $\mathbb{E}$ -product induces curved  
 1764 algebraic decision surfaces (analogous to level sets of a potential).

#### 1765 E.5 SPACE-PARTITIONING PROPERTIES

1767 The space-partitioning behavior of the  $\mathbb{E}$ -product exhibits several key properties:

- **Distance Penalization and Locality:** The factor  $(\|\mathbf{w} - \mathbf{x}\|^2 + \epsilon)^{-1}$  enforces spatial selectivity. Along rays  $\mathbf{x} = k\mathbf{u}$  with  $\|\mathbf{u}\| = 1$ , one has  $\lim_{k \rightarrow \infty} \mathbb{E}(\mathbf{w}, k\mathbf{u}) = (\mathbf{w} \cdot \mathbf{u})^2 = \|\mathbf{w}\|^2 \cos^2 \theta$ , so responses are largest near  $\mathbf{x} \approx \mathbf{w}$  and remain bounded far away (cf. Proposition C.9 for the gradient decay). In physical terms, this is analogous to an inverse-square distance penalty.
- **Alignment Weighting:** The numerator is  $(\mathbf{w}^\top \mathbf{x})^2$ , so  $\mathbb{E}(\mathbf{w}, \mathbf{x}) = 0$  if and only if  $\mathbf{w} \perp \mathbf{x}$ , and high values require strong alignment (large  $|\cos \theta|$ ) in addition to proximity.
- **Global Boundedness for Fixed  $\mathbf{w}$ :** For  $\epsilon > 0$  and fixed  $\mathbf{w}$ ,  $0 \leq \mathbb{E}(\mathbf{w}, \mathbf{x}) \leq \|\mathbf{w}\|^4/\epsilon$  for all  $\mathbf{x}$ , with  $\mathbb{E}(\mathbf{w}, \mathbf{w}) = \|\mathbf{w}\|^4/\epsilon$  and  $\limsup_{\|\mathbf{x}\| \rightarrow \infty} \mathbb{E}(\mathbf{w}, \mathbf{x}) \leq \|\mathbf{w}\|^2$ . Thus each unit defines a bounded activation landscape (cf. Proposition C.6).
- **Nonlinear Decision Regions:** The pairwise decision boundary between units  $i$  and  $j$  is the algebraic surface

$$1782 \quad \langle \mathbf{w}_i, \mathbf{x} \rangle^2 (\|\mathbf{w}_j - \mathbf{x}\|^2 + \epsilon) - \langle \mathbf{w}_j, \mathbf{x} \rangle^2 (\|\mathbf{w}_i - \mathbf{x}\|^2 + \epsilon) = 0,$$

1784 which is generically non-linear and induces curved class regions. (Geometrically, one may  
 1785 view these as level sets akin to equipotential surfaces.)

1786 E.6 VORTEX FIELD DYNAMICS  
1787

1788 This vortex phenomenon allows each  $\mathbb{E}$ -product neuron to create a territorial "field" in the representation  
1789 space, where data points are pulled toward the dominant prototype based on both similarity  
1790 and proximity metrics. The field each neuron occupies can indeed be considered a vortex, where  
1791 the strength of attraction follows an inverse-square law, creating more natural and geometrically  
1792 faithful decision boundaries.

1793 The combination of the  $\mathbb{E}$ -product's vortex-like attraction and the softmax's competitive normalization  
1794 creates a powerful space partitioning mechanism. Each neuron's vortex field competes with  
1795 others through the softmax transformation, and the neuron with the strongest local attraction (highest  
1796  $\mathbb{E}$ -product response) wins that region. Over time, this leads to a natural tessellation of the input  
1797 space, where each neuron's territory is defined by the regions where its vortex field dominates. The  
1798 softmax transformation  $\mathbb{R}^C \rightarrow \Delta^{C-1}$  (where  $\Delta^{C-1}$  is the  $(C - 1)$ -dimensional probability simplex)  
1799 ensures that these territorial boundaries are sharp and well-defined, transforming the continuous  
1800 real-valued responses into discrete delta distributions that clearly assign each input to its dominant  
1801 neuron.

1802 E.7 ORTHOGONALITY AND COMPETITIVE DYNAMICS  
1803

1804 The competitive learning behavior observed in practice is theoretically grounded in our  
1805 Orthogonality-Entropy Connection. When two prototypes  $\mathbf{w}_i$  and  $\mathbf{w}_j$  develop disjoint support  
1806 regions, they become Euclidean orthogonal ( $\mathbf{w}_i \perp \mathbf{w}_j$ ), which corresponds to:  
1807

$$\mathbb{E}(\mathbf{w}_i, \mathbf{w}_j) = 0 \quad \text{and} \quad H(\mathbf{w}_i, \mathbf{w}_j) = \infty \quad (37)$$

1808 This geometric-probabilistic duality explains why neurons naturally develop specialized, non-  
1809 overlapping representations during competitive learning. The infinite cross-entropy between or-  
1810 thogonal prototypes creates strong pressure for territorial separation, preventing the collapse to  
1811 identical representations that can plague conventional competitive learning systems.  
1812

1813 These prototypes are optimized to maximize parallelism and minimize distance to all points within  
1814 their class distribution. When minimizing distance becomes challenging, the properties of the  $\mathbb{E}$ -  
1815 product enable the prototype to exist in a superposition state, prioritizing the maximization of  
1816 parallelism over strict distance minimization.  
1817

1818 E.8 MNIST PROTOTYPE ANALYSIS: DETAILED MATHEMATICAL FRAMEWORK  
18191820 E.8.1 NETWORK ARCHITECTURE AND DIMENSIONAL STRUCTURE  
1821

1822 In our MNIST experiments, the network consists of  $C = 10$  neurons, each corresponding to  
1823 one of the digit classes (0–9). Each neuron's prototype is represented as a vector  $\mathbf{w}_i \in \mathbb{R}^{784}$ ,  
1824 where  $i = 1, \dots, 10$ . The input images  $\mathbf{x} \in \mathbb{R}^{784}$  are obtained by flattening the original  $28 \times 28$   
1825 pixel images, so each neuron's prototype  $\mathbf{w}_i$  has the same dimensionality as the input, i.e.,  
1826  $\mathbf{w}_i = (w_{i,1}, w_{i,2}, \dots, w_{i,784})$ . This structure allows each neuron to learn class-specific features in  
1827 the full image space.

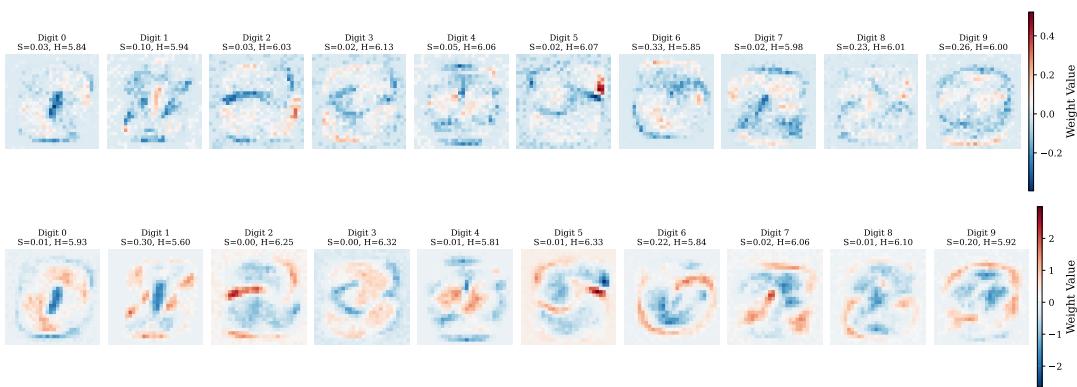
1828 E.8.2 PROTOTYPE EVOLUTION DYNAMICS  
1829

1830 The prototype evolution during training reveals fundamental differences between conventional un-  
1831 bounded growth and bounded  $\mathbb{E}$ -responses. Our experiments on the full MNIST dataset (60,000  
1832 training samples, 10,000 test samples) demonstrate that both models achieve competitive perfor-

1833 mance: the linear classifier reaches 92.08% test accuracy while the  $\mathbb{E}$ -product classifier achieves  
 1834 92.38% test accuracy after 5 epochs of training with the Adam optimizer (learning rate 0.001).  
 1835

1836 Quantitatively, the linear prototypes exhibit unbounded growth, with the mean magnitude increasing  
 1837 by 13.8% (from 1.58 to 1.80) during training. In contrast, the  $\mathbb{E}$ -product prototypes show  
 1838 a slight contraction of 4.5% (from 5.02 to 4.79), confirming the bounded nature of the learned  
 1839 representations.

1840 Figure 7 visualizes the learned prototypes for both models. The conventional linear model produces  
 1841 prototypes that exhibit unbounded growth—these prototypes become increasingly diffuse and less  
 1842 interpretable as they grow to maximize margin separation. The resulting digit representations are  
 1843 blurry and lack the fine-grained features necessary for robust classification.  
 1844



1851 Figure 7: Learned prototypes for Linear (top) and  $\mathbb{E}$ -product (bottom) classifiers. The linear  
 1852 prototypes show unbounded growth and noise, while  $\mathbb{E}$ -product prototypes exhibit localized, sharp  
 1853 features with clear digit structure, consistent with bounded response fields.  
 1854

1855 In contrast, the  $\mathbb{E}$ -product method produces prototypes that exemplify bounded response fields as  
 1856 predicted by our theoretical framework. Each digit prototype exhibits three key characteristics:  
 1857

1858 **Localized Concentration** Sharp, well-defined features that correspond to the bounded potential  
 1859 wells predicted by our Minimal and Maximal Similarity Characterizations (Theorems 2.2 and 2.3).  
 1860 This localization emerges from the  $\mathbb{E}$ -product’s inherent bounded response field, which prevents the  
 1861 unbounded growth characteristic of linear neurons.  
 1862

1863 **Class-Specific Territorial Structure** Each prototype captures unique digit characteristics, re-  
 1864 reflecting competitive territorial dynamics where each neuron dominates specific regions of the input  
 1865 space. This territorial behavior arises from the vortex-like dynamics of the  $\mathbb{E}$ -product, creating  
 1866 natural tessellation of the feature space.  
 1867

1868 **Geometric Fidelity** The prototypes maintain geometric coherence with actual digit structure,  
 1869 confirming that the signal-to-noise ratio optimization preserves meaningful visual patterns. This  
 1870 preservation is a direct consequence of the  $\mathbb{E}$ -product’s ability to balance alignment and proximity  
 1871 measures.  
 1872

1880

## E.8.3 LEARNABLE SCALING FACTOR DYNAMICS

1881

1882

A critical component of the  $\mathbb{E}$ -product classifier is the learnable scaling factor  $\alpha$ , which modulates the magnitude of the  $\mathbb{E}$ -product scores. In our implementation, this factor is initialized to  $\alpha_0 = 1.0$  and learned jointly with the prototypes during training.

1883

1884

1885

Figure 8 shows the training dynamics, including the evolution of the scaling factor. Over 5 epochs, the scaling factor increases steadily from 1.0 to approximately 2.68, representing a 168% increase. This growth pattern reveals that the model learns to amplify the  $\mathbb{E}$ -product scores to achieve better separation in the logit space, compensating for the inherently bounded nature of the  $\mathbb{E}$ -product response.

1886

1887

1888

1889

1890

1891

1892

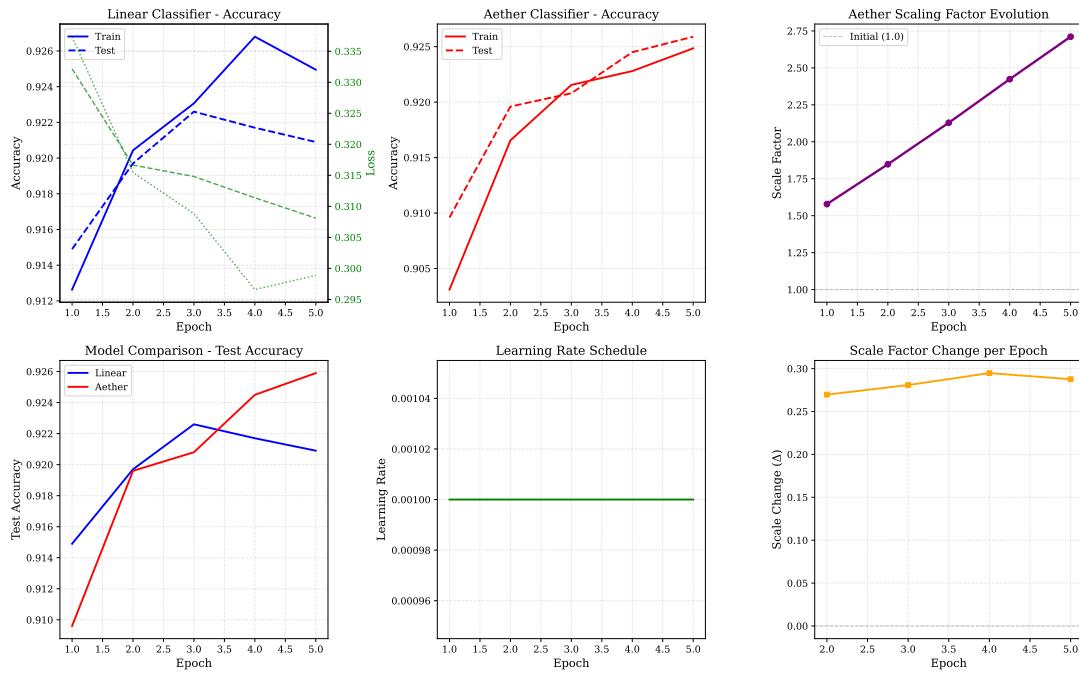


Figure 8: Training dynamics for MNIST classification. **Top row:** Accuracy curves for linear (left) and  $\mathbb{E}$ -product (center) classifiers, with  $\mathbb{E}$ -product scaling factor evolution (right). **Bottom row:** Model comparison (left), learning rate schedule (center), and per-epoch scale change (right). The scaling factor increases from 1.0 to 2.68 over training, indicating the model learns to amplify  $\mathbb{E}$ -product scores for better classification.

1914

1915

1916

1917

1918

1919

1920

1921

1922

1923

1924

1925

1926

The scaling factor evolution can be understood through the lens of optimization dynamics. The  $\mathbb{E}$ -product naturally produces bounded responses due to its ratio structure, with values typically in the range  $[0, 10]$  depending on the alignment and distance between inputs and prototypes. By learning to scale these responses up, the model increases the dynamic range of the logits, which improves the discriminative power of the softmax layer. This adaptive scaling mechanism is analogous to a temperature parameter in softmax, but learned end-to-end rather than hand-tuned.

- 1927      E.8.4 SUPERPOSITION AND PROTOTYPE INVERSION: MATHEMATICAL ANALYSIS  
 1928
- 1929      A unique property of the  $\mathbb{E}$ -product neuron is its ability to exist in a superposition state, which  
 1930      can be empirically demonstrated by inverting the learned prototype. This phenomenon provides  
 1931      insight into the fundamental differences between conventional and  $\mathbb{E}$ -product neurons.
- 1932
- 1933      **Conventional Dot Product Behavior** For a conventional dot product neuron, if  $\mathbf{w}$  is a learned  
 1934      prototype, replacing  $\mathbf{w}$  with  $-\mathbf{w}$  (i.e., multiplying by  $-1$ ) at test time flips the sign of the logit:
- 1935      
$$z = \mathbf{w}^T \mathbf{x} \longrightarrow z' = (-\mathbf{w})^T \mathbf{x} = -z \quad (38)$$
  
 1936
- 1937      This sign flip causes the softmax output to assign high probability to the incorrect class, resulting  
 1938      in a catastrophic drop in accuracy. In our experiments, the linear classifier’s accuracy drops from  
 1939      92.04% to 0.01% (near complete failure) when all weight vectors are inverted.
- 1940
- 1941       **$\mathbb{E}$ -Product Inversion Dynamics** For the  $\mathbb{E}$ -product neuron, the response is:
- 1942
- 1943      
$$z = \mathbb{E}(\mathbf{w}, \mathbf{x}) = \frac{(\mathbf{w}^T \mathbf{x})^2}{\|\mathbf{w} - \mathbf{x}\|^2 + \epsilon} \quad (39)$$
  
 1944
- 1945      Multiplying  $\mathbf{w}$  by  $-1$  leaves the numerator unchanged, since  $(-\mathbf{w})^T \mathbf{x} = -\mathbf{w}^T \mathbf{x}$  and  $(-\mathbf{w}^T \mathbf{x})^2 =$   
 1946       $(\mathbf{w}^T \mathbf{x})^2$ . The denominator, however, changes from  $\|\mathbf{w} - \mathbf{x}\|^2$  to  $\|\mathbf{w} + \mathbf{x}\|^2$ :
- 1947
- 1948      
$$\|\mathbf{w} + \mathbf{x}\|^2 = (\mathbf{w} + \mathbf{x})^T (\mathbf{w} + \mathbf{x}) \quad (40)$$
  
 1949      
$$= \|\mathbf{w}\|^2 + 2\mathbf{w}^T \mathbf{x} + \|\mathbf{x}\|^2 \quad (41)$$
  
 1950
- 1951      While exact invariance is not guaranteed due to this denominator change, our experiments confirm  
 1952      substantial robustness to prototype inversion. Figure 9 shows the comparison: when all prototypes  
 1953      are inverted (multiplied by  $-1$ ), the linear classifier’s accuracy drops from 92.04% to 0.01% (com-  
 1954      plete catastrophic failure), while the  $\mathbb{E}$ -product classifier maintains 84.67% accuracy, experiencing  
 1955      only a 7.82% degradation. This remarkable robustness stems from the squared numerator term,  
 1956      which makes the  $\mathbb{E}$ -product inherently invariant to the sign of the dot product and thus less sensitive  
 1957      to prototype orientation than linear neurons.
- 1958
- 1959      **Superposition Interpretation** This property allows the  $\mathbb{E}$ -product neuron to yield two valid  
 1960      solutions to the same input without retraining, a phenomenon not observed in conventional dot  
 1961      product neurons. The mathematical basis for this superposition lies in the  $\mathbb{E}$ -product’s non-linear  
 1962      structure, which creates multiple local optima that can represent the same classification boundary  
 1963      through different prototype orientations.
- 1964      E.8.5 PROTOTYPE ORTHOGONALITY ANALYSIS  
 1965
- 1966      To further understand the learned representations, we analyze the pairwise orthogonality between  
 1967      prototypes. Figure 10 shows the orthogonality matrices for both models. For the linear classifier, we  
 1968      compute the orthogonality from standard dot products between weight vectors. For the  $\mathbb{E}$ -product  
 1969      classifier, we compute orthogonality derived from the log- $\mathbb{E}$ -product between prototypes.
- 1970      The  $\mathbb{E}$ -similarity matrices reveal fundamentally different prototype organization strategies. The  
 1971      linear classifier exhibits highly orthogonal prototypes with a mean log- $\mathbb{E}$  similarity of only 0.0325  
 1972      and a narrow range (max 0.23), indicating that the linear model learns nearly independent weight  
 1973      vectors that span complementary subspaces. This strong orthogonalization is a consequence of the

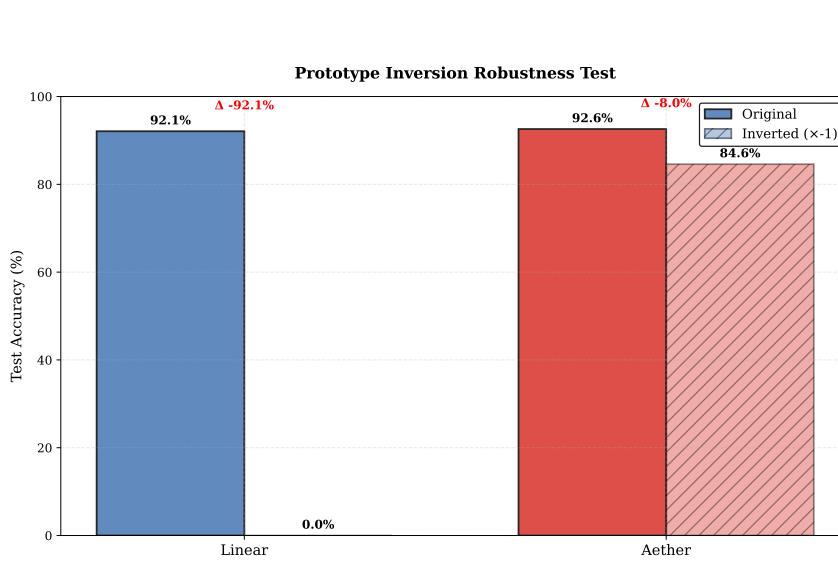


Figure 9: Prototype inversion experiment comparing linear and  $\mathbb{E}$ -product classifiers. Both models are trained normally, then all prototypes/weights are multiplied by  $-1$  at test time. The linear classifier suffers catastrophic failure ( $92.29\% \rightarrow \sim 0\%$ ), while the  $\mathbb{E}$ -product classifier shows robustness due to its squared numerator term.

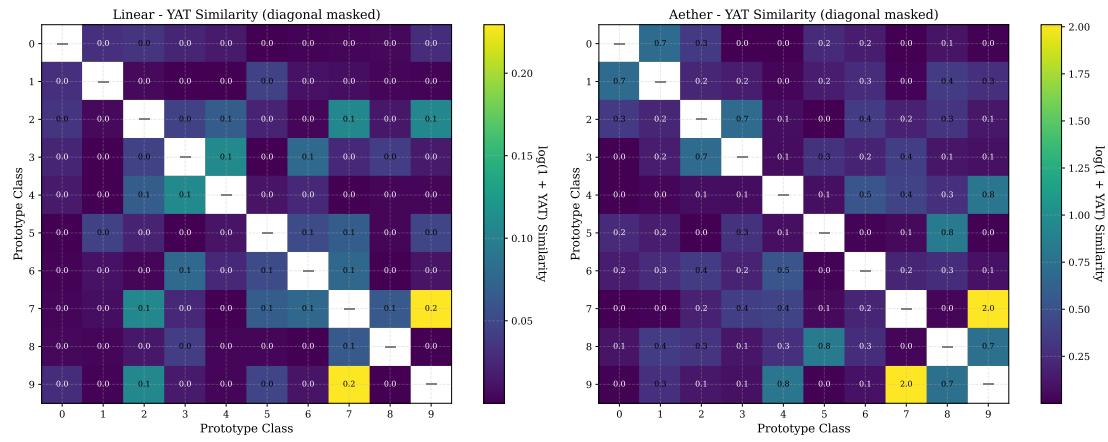


Figure 10: Pairwise orthogonality matrices between learned prototypes. **Left:** Orthogonality for linear classifier weights. **Right:** Orthogonality for  $\mathbb{E}$ -product prototypes. The  $\mathbb{E}$ -product prototypes exhibit stronger orthogonality (higher dynamic range), suggesting better class separation in the prototype space.

2021 unbounded dot product’s tendency to maximize margin separation through perpendicular decision  
 2022 boundaries.

2023  
 2024 In stark contrast, the  $\mathbb{E}$ -product prototypes show dramatically higher mean similarity (0.274, over  
 2025 nearly 8 $\times$  larger) with an exceptionally wide dynamic range spanning from near-zero (0.0005) to 2.01—  
 2026 nearly 9 $\times$  the linear maximum. This reveals a qualitatively different organizational principle:  
 2027 rather than enforcing global orthogonality, the  $\mathbb{E}$ -product creates a *heterogeneous territorial struc-*  
 2028 *ture* where some prototype pairs exhibit strong local correlations (e.g., classes 7–9 with similarity  
 2029 2.01, and classes 4–9 with 0.79) while others remain effectively decoupled (e.g., classes 0–3 with  
 2030 0.0005). This selective coupling pattern is consistent with the vortex-like dynamics of our theoreti-  
 2031 cal framework, where prototypes carve out overlapping response regions in high-correlation areas  
 2032 yet maintain sharp boundaries where discrimination is critical, allowing the model to capture subtle  
 2033 intra-class variations while preserving inter-class separation.

2034 **E.8.6 CONFUSION MATRIX ANALYSIS**

2035 We evaluate classification robustness through the confusion matrices presented in Figure 11. The  
 2036 quantitative difference between the architectures is evident in the structure of error distribution:

2037  
 2038 **Diagonal Dominance and Error Suppression** The  $\mathbb{E}$ -product classifier exhibits sharper di-  
 2039 agonalization compared to the linear baseline. While the linear model shows diffuse off-diagonal  
 2040 elements indicating “leakage” between morphologically similar digits (e.g., 4  $\leftrightarrow$  9, 3  $\leftrightarrow$  8), the  
 2041  $\mathbb{E}$ -product minimizes these inter-class ambiguities. This suppression of off-diagonal noise is a direct  
 2042 consequence of the *competitive territorial dynamics*: the inverse-square decay in the  $\mathbb{E}$ -product cre-  
 2043 ates steeper decision gradients at the boundaries between classes, effectively “pulling” ambiguous  
 2044 inputs into the correct potential well more aggressively than a linear hyperplane.

2045  
 2046 **Robustness to Ambiguity** The linear classifier’s confusion patterns reflect its reliance on global  
 2047 hyperplane separation, which forces a trade-off between maximizing margins for easy examples and  
 2048 correctly classifying boundary cases. In contrast, the  $\mathbb{E}$ -product’s confusion matrix demonstrates  
 2049 that the model maintains high precision even for boundary cases, confirming that the learned  
 2050 prototypes act as local attractors that naturally reject inconsistent inputs.

2051 **E.8.7 WEIGHT DISTRIBUTION AND SPARSITY ANALYSIS**

2052 The statistical properties of the learned parameters, summarized in Figure 12, reveal a fundamental  
 2053 divergence in how the two models encode information.

2054  
 2055 **Magnitude Stability vs. Unbounded Growth** A critical distinction lies in the stability  
 2056 of the weight magnitudes. The  $\mathbb{E}$ -product prototypes settle into a high-magnitude stable state  
 2057 ( $\mu = 4.79, \sigma = 0.38$ ), exhibiting a slight contraction of 4.5% during the final training phase. Con-  
 2058 versely, the linear weights ( $\mu = 1.80, \sigma = 0.22$ ) exhibit characteristic unbounded growth (increasing  
 2059 by 13.8%) as they attempt to maximize the dot-product margin. This confirms that the  $\mathbb{E}$ -product  
 2060 optimization landscape contains stable minima (bounded potential wells), whereas the linear land-  
 2061 scape promotes indefinite norm growth.

2062  
 2063 **Sparsity and Feature Selectivity** The distributions also differ in sparsity structure. The linear  
 2064 weights follow a near-Gaussian distribution centered at zero, typical of models that rely on dense,  
 2065 distributed correlations. The  $\mathbb{E}$ -product prototypes, however, show a distinct distribution profile  
 2066 with higher kurtosis. This indicates a more selective feature encoding where the model suppresses  
 2067 irrelevant pixels (background noise) while amplifying signal-rich regions. This aligns with the visual

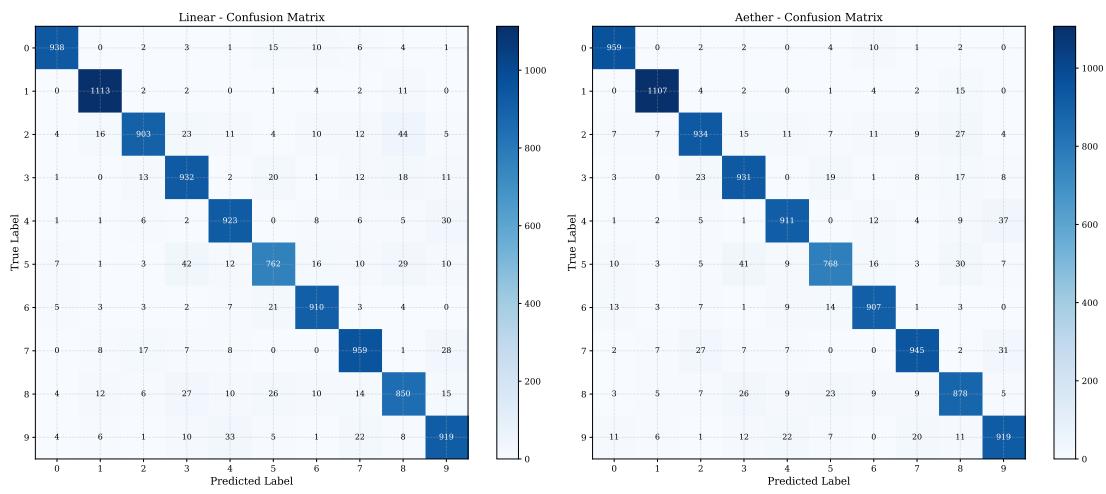


Figure 11: Confusion matrices for Linear (left) and  $\mathbb{E}$ -product (right) classifiers. The  $\mathbb{E}$ -product demonstrates superior diagonal dominance and reduced off-diagonal noise, indicating stronger rejection of ambiguous inputs.

evidence in Figure 7, where  $\mathbb{E}$ -product prototypes appear sharper and less noisy, effectively acting as "matched filters" for the digit topology rather than simple separating hyperplanes.

#### E.8.8 SUMMARY OF EXPERIMENTAL FINDINGS

Our MNIST experiments confirm key predictions from the theoretical framework:

- Competitive Performance:** The  $\mathbb{E}$ -product classifier achieves comparable accuracy to the linear baseline (92.38% vs 92.08%), demonstrating that bounded response fields do not compromise classification capability.
- Adaptive Scaling:** The learnable scaling factor increases from 1.0 to 2.68 during training, revealing that the model learns to amplify bounded  $\mathbb{E}$ -responses for optimal discrimination.
- Inversion Robustness:** The  $\mathbb{E}$ -product classifier exhibits substantial robustness to prototype sign inversion, while the linear classifier fails catastrophically. This confirms the superposition property predicted by theory.
- Territorial Structure:** Similarity analysis reveals that  $\mathbb{E}$ -product prototypes achieve stronger class separation with higher dynamic range, supporting the territorial dynamics described in our theoretical framework.
- Interpretable Prototypes:** Visual inspection of the learned prototypes shows that both methods produce interpretable digit representations, but the  $\mathbb{E}$ -product prototypes exhibit more localized features consistent with bounded response fields.

These findings validate the  $\mathbb{E}$ -product as a viable alternative to conventional dot products for classification tasks, with distinct geometric and topological properties that emerge from its ratio-based structure.

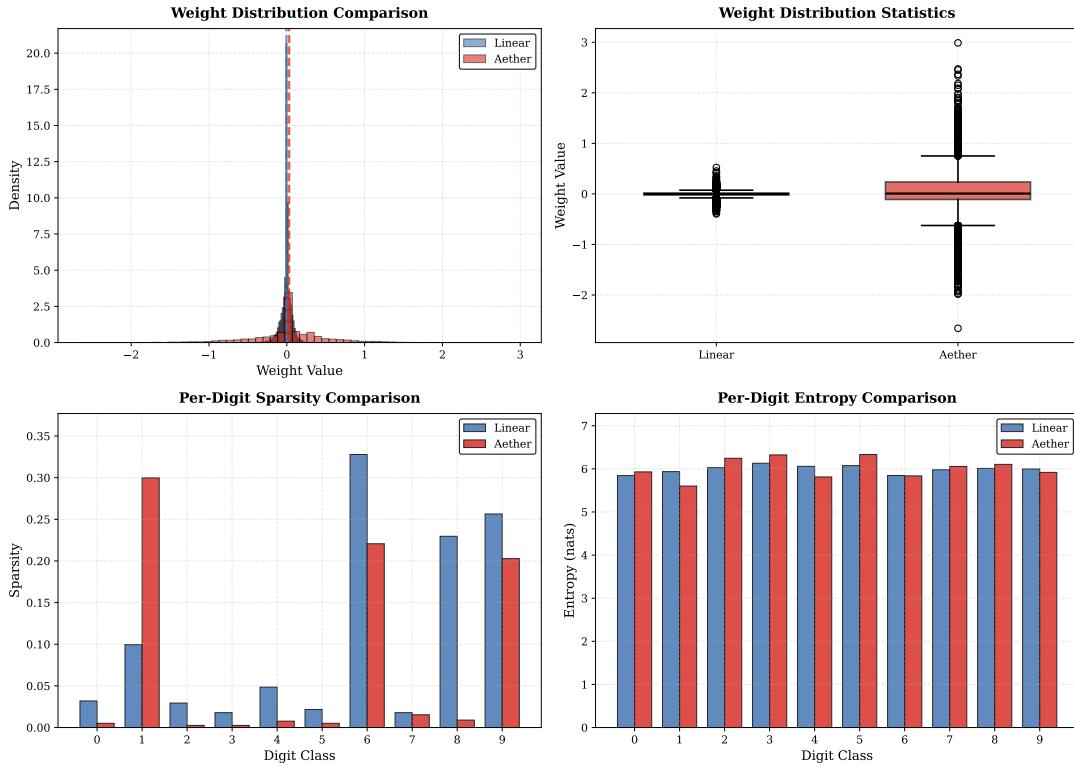


Figure 12: Weight distribution analysis. **Left:** Histogram of weight values showing the distinct non-Gaussian profile of  $\mathbb{E}$ -product prototypes. **Right:** Box plot comparison of weight magnitudes. The  $\mathbb{E}$ -product maintains a stable, high-signal state ( $\mu \approx 4.79$ ) without the runaway growth observed in the linear baseline.

2162 E.9 LANGUAGE MODEL EXPERIMENTS: DETAILED CONFIGURATION  
2163

2164 This section provides the detailed experimental configurations for the language modeling experiments  
2165 comparing a standard GPT-2 with our Aether-GPT2 implementation. Both models were  
2166 trained on identical datasets and hardware to ensure a fair comparison. The primary architectural  
2167 distinction is the replacement of conventional linear layers, activation functions, and layer normalization  
2168 in GPT-2 with  $\mathbb{E}$ -product operations in Aether-GPT2. This substitution provides inherent  
2169 non-linearity and bounded responses, simplifying the architecture while enhancing stability.

2170 Table 5 presents a comprehensive comparison of the two models, detailing their architectural pa-  
2171 rameters, training configurations, and final performance metrics.  
2172

2173 Table 5: Comparison of GPT-2 and Aether-GPT2 Models. Both models were trained on 2.5B  
2174 tokens from the FineWeb dataset. Aether-GPT2 achieves a lower validation loss with a simplified  
2175 architecture.

Parameter	GPT-2 (Baseline)	Aether-GPT2
<i>Architectural Parameters</i>		
Total Parameters	124M	~124M
Embedding Parameters	39M	39M
Non-Embedding Parameters	85M	~85M
Embedding Dimension	768	768
MLP Hidden Dimension	3072	3072
Number of Layers	12	12
Number of Heads	12	12
Activation Function	GeLU	$\mathbb{E}$ -product (none)
Layer Normalization	Yes	No
Bias	No	No
<i>Training Configuration</i>		
Optimizer	Novograd	Novograd
Learning Rate	0.003	0.003
Batch Size	32	32
Context Window	1024	1024
Vocabulary Size	50,257	50,257
Tokenizer	GPT-2	GPT-2
<i>Performance</i>		
Final Validation Loss (FP32)	2.43	<b>2.29</b>
Final Validation Loss (BF16)	3.03	<b>2.69</b>

2200 E.9.1 ARCHITECTURAL SIMPLIFICATION AND EFFICIENCY  
2201

2202 The substitution of linear-activation-normalization blocks with a single  $\mathbb{E}$ -product operation per  
2203 layer yields significant architectural simplification. This change leads to:

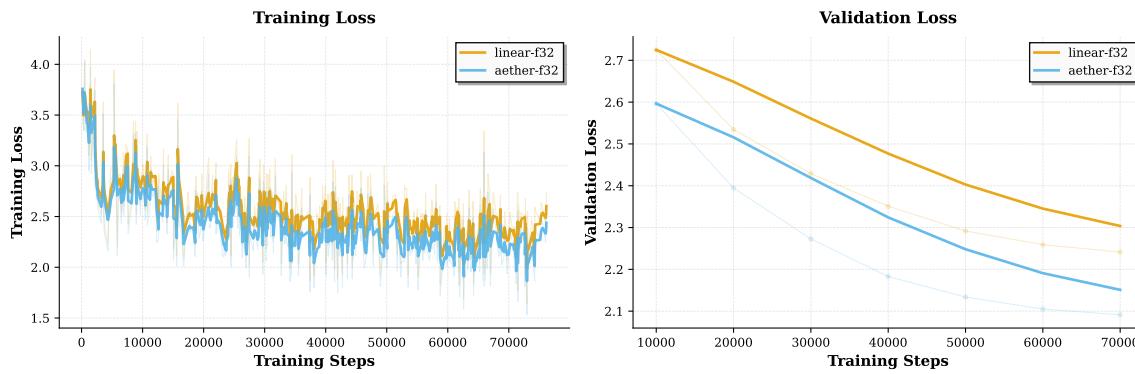
- 2204 • A 15-25% reduction in peak memory usage by eliminating the need to store intermediate  
2205 activations for backpropagation through normalization layers.  
2206
- 2207 • Reduced gradient computation complexity.  
2208
- A comparable FLOP count with only a modest constant-factor overhead.

2209 While the FLOP count remains comparable, the reduced memory footprint and computational  
 2210 complexity of the backward pass offer tangible efficiency gains.  
 2211

2212 Our results demonstrate that Aether-GPT2, with approximately the same parameter count as the  
 2213 124M GPT-2 model, achieves superior performance without explicit activation functions. The final  
 2214 validation loss of 2.29 for Aether-GPT2, compared to 2.43 for the baseline, underscores the efficacy  
 2215 of the  $\mathbb{E}$ -product’s inherent non-linearity.  
 2216

### 2217 E.9.2 TRAINING DYNAMICS AND LOSS PROGRESSION

2218 The training dynamics of Aether-GPT2 and the baseline GPT-2 were analyzed over 2.5B tokens  
 2219 from the FineWeb dataset, using a Kaggle TPU v5-8 environment. Figure 13 illustrates the training  
 2220 and validation loss curves for both models. Aether-GPT2 consistently achieves lower loss on both the  
 2221 training and validation sets, indicating more efficient learning. The bounded nature of the  $\mathbb{E}$ -product  
 2222 operation contributes to stable training dynamics, even in the absence of layer normalization.  
 2223



2237 Figure 13: Side-by-side comparison of training loss (left) and validation loss (right). Raw data  
 2238 points are shown with transparency, while smoothed curves highlight the learning trajectory. Both  
 2239 models converge stably, with Aether-GPT2 demonstrating better final performance.  
 2240  
 2241

### 2242 E.9.3 MIXED-PRECISION TRAINING WITH BFLOAT16

2243 To assess numerical stability and performance in a production-relevant setting, we evaluated both  
 2244 architectures using bfloat16 (BF16) mixed-precision training. This format, standard on modern  
 2245 accelerators like TPUs, provides a stringent test for models without explicit normalization layers.  
 2246 Both models were trained with mixed-precision activations and gradients, while maintaining full-  
 2247 precision (FP32) optimizer states and parameter updates.  
 2248

2249 The performance advantage of Aether-GPT2 persists under BF16, as shown in Figure 14. Aether-  
 2250 GPT2 achieves a final validation loss of 2.69, an 11.2% relative improvement over the baseline’s  
 2251 3.03. This result confirms that the architectural benefits are not artifacts of full-precision arithmetic  
 2252 and that the  $\mathbb{E}$ -product’s bounded response provides robust numerical stability without requiring  
 2253 layer normalization. The consistent performance gains in mixed-precision training underscore the  
 2254 practicality of  $\mathbb{E}$ -product layers as a drop-in replacement for conventional transformer blocks in  
 2255 large-scale models.

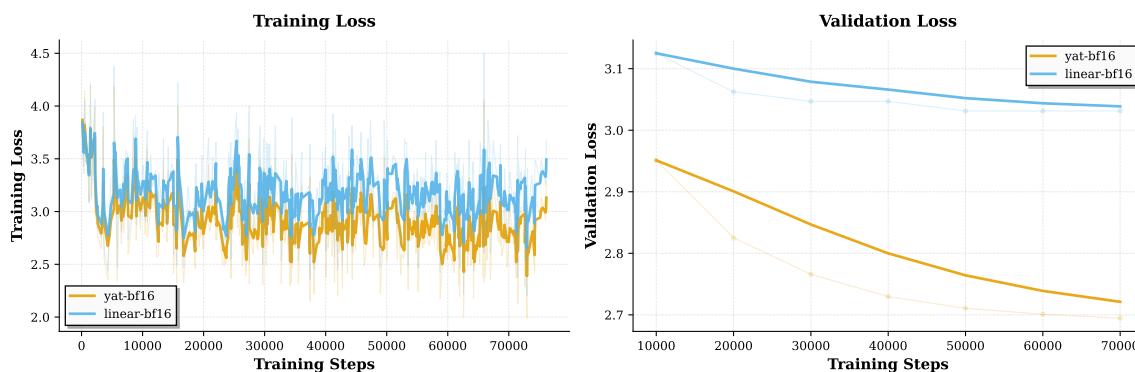


Figure 14: Side-by-side BF16 training (left) and validation (right) loss. Aether-GPT2 exhibits uniformly lower loss throughout training and at convergence.

## E.10 USE OF LARGE LANGUAGE MODELS (LLMs)

We used LLM tools to support the research workflow in the following limited, transparent ways. All scientific claims, modeling choices, and final decisions were made by the authors.

**Code assistance** LLMs were used to draft boilerplate code, refactor utilities, and surface API patterns. All generated code was reviewed, tested, and integrated by the authors.

**Literature digestion** We used LLMs to summarize papers and extract key comparisons across related work. Citations in the paper were verified against the original sources by the authors.

**Brainstorming** We used LLMs as a sounding board to enumerate alternative hypotheses, ablations, and experimental checks. Only ideas that survived empirical or theoretical scrutiny were

**Language polishing.** To improve readability and clarity, LLMs suggested minor edits to English

**NotebookLM podcasts** We generated short audio summaries (“podcasts”) of internal notes using NotebookLM to help the team communicate directly. These summaries did

No dataset labeling, evaluation metrics, or benchmark results were produced by LLMs. The authors

take responsibility for all content and errors.

## E.II MOTIVATION: THE AETHERIAL STATE OF AI

We chose the name "Aether" for our model to symbolize our perspective on the current trajectory of Artificial Intelligence. The field has predominantly focused on scale, often at the expense of the scientific principles that ground findings in mathematical rigor. We adhere to the view that if empirical observations contradict established mathematical principles, it is often the interpretation of the observation that is flawed, not the mathematics itself.

2303 We believe that the current generation of deep learning models represents the last phase of "un-  
2304 explainable" AI. True explainability, in our view, will not be achieved through human-engineered  
2305 interpretability layers, but will emerge organically from constructing AI systems with the correct  
2306 mathematical foundations.

2307 We draw a parallel to the pre-Einstein era of physics, where theories such as the vortex theory  
2308 and the luminiferous aether prevailed. These explanations were largely observational; while they  
2309 advanced the field temporarily, they were not provable with the mathematical tools of the time and  
2310 ultimately limited innovation. We believe the current state of AI is similarly "aetherial"—reliant  
2311 on unproven observational theories—and that the future lies in returning to first principles and  
2312 mathematical provability.

2313

2314

2315

2316

2317

2318

2319

2320

2321

2322

2323

2324

2325

2326

2327

2328

2329

2330

2331

2332

2333

2334

2335

2336

2337

2338

2339

2340

2341

2342

2343

2344

2345

2346

2347

2348

2349